



中国科学技术大学
University of Science and Technology of China

Image Convolution

Juyong Zhang
School of Mathematics, USTC

Spatial Filtering

Let \mathbf{I} and \mathbf{J} be images such that $\mathbf{J} = T[\mathbf{I}]$.

$T[\cdot]$ represents a transformation, such that,

$$\begin{aligned}\mathbf{J}(r,c) &= T[\mathbf{I}](r,c) = \\ f\left(\left\{\mathbf{I}(\rho, \chi) \mid \rho \in \{r-s, \dots, r, \dots, r+s\}, \chi \in \{c-d, \dots, c, \dots, c+d\}\right\}\right).\end{aligned}$$

That is, the value of the transformed image, \mathbf{J} , at pixel location (r,c) is a function of the values of the original image, \mathbf{I} , in a $2s+1 \times 2d+1$ rectangular neighborhood centered on pixel location (r,c) .

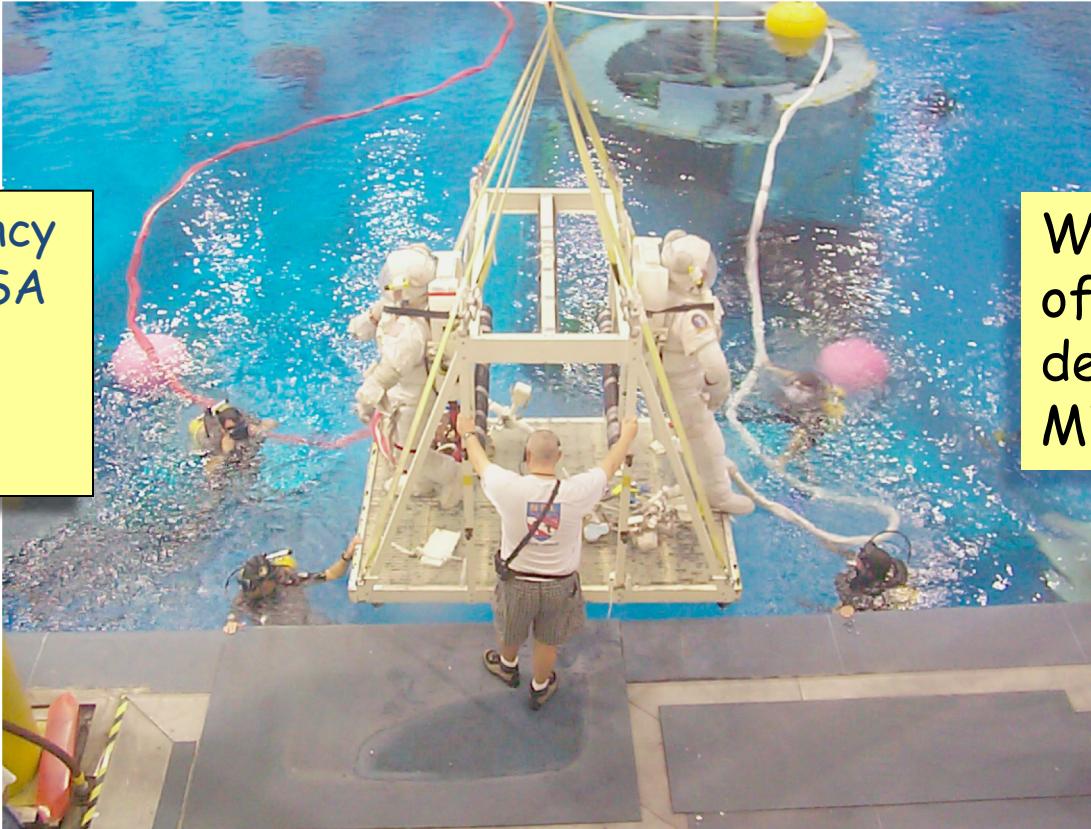


Moving Windows

- The value, $J(r,c) = T[I](r,c)$, is a function of a rectangular neighborhood centered on pixel location (r,c) in I .
- There is a different neighborhood for each pixel location, but if the dimensions of the neighborhood are the same for each location, then transform T is sometimes called a moving window transform.



Moving-Window Transformations

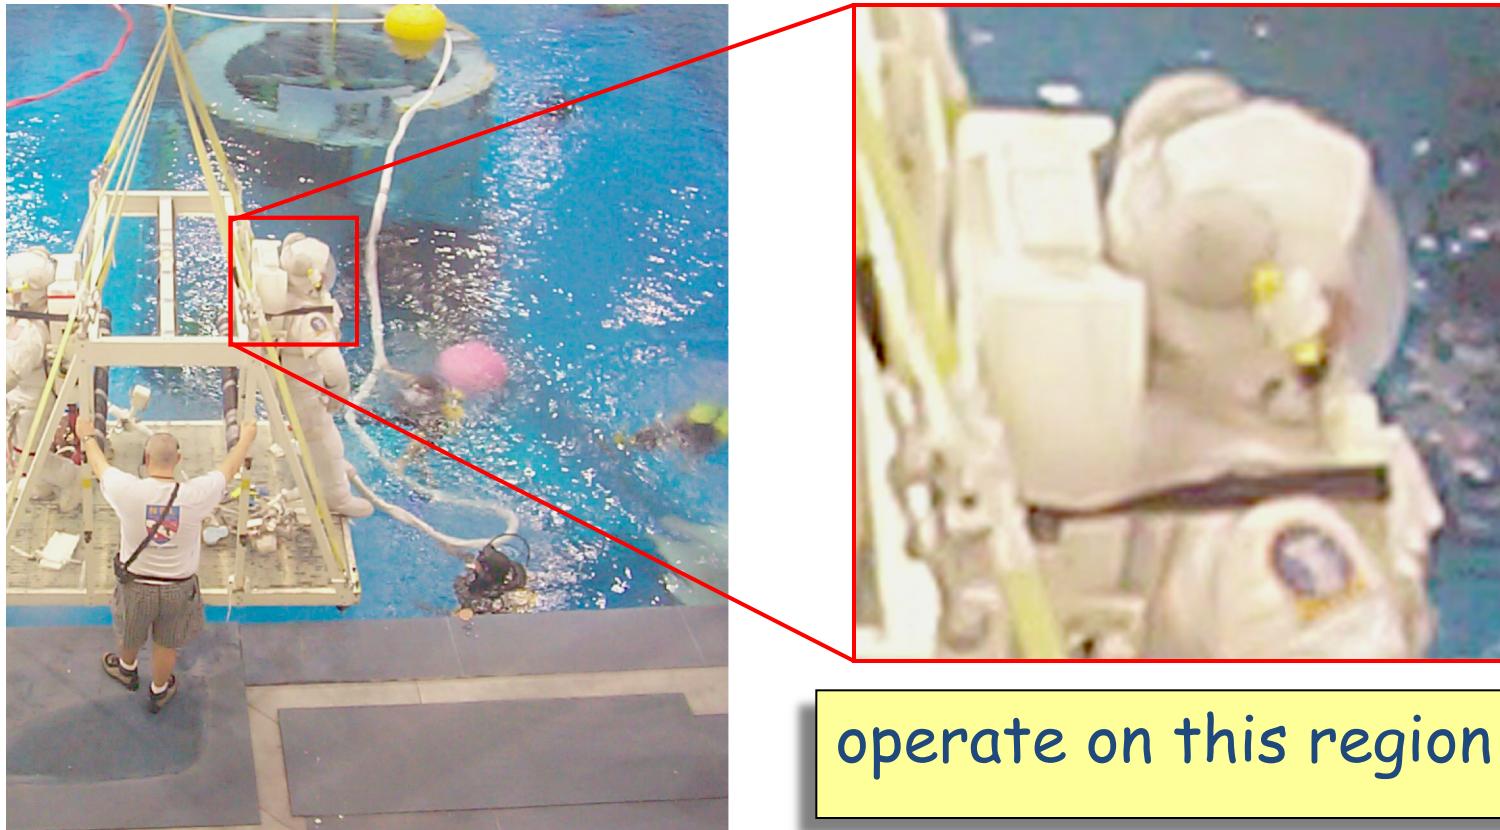


Neutral Buoyancy Facility at NASA Johnson Space Center

We'll take a section of this image to demonstrate the MWT

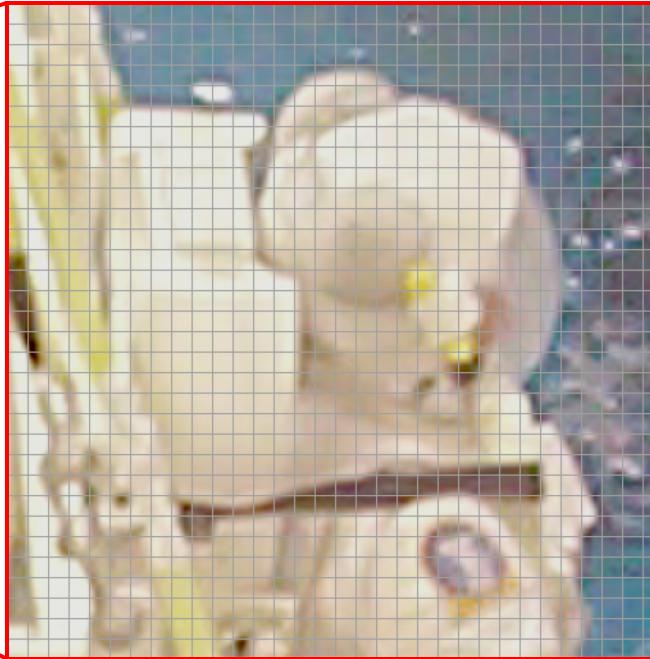
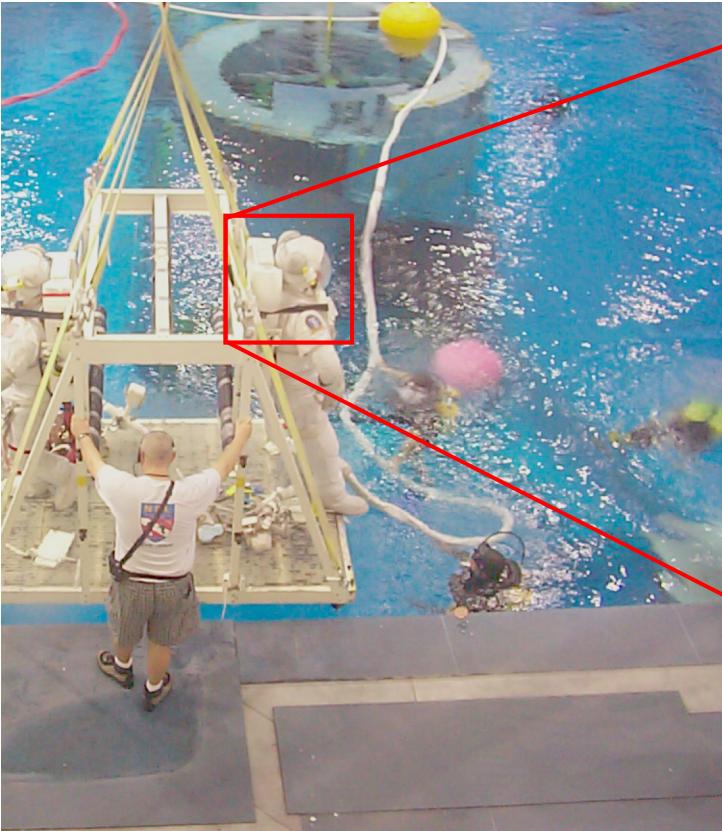
photo: R.A.Peters II, 1999

Moving-Window Transformations



Moving-Window Transformations

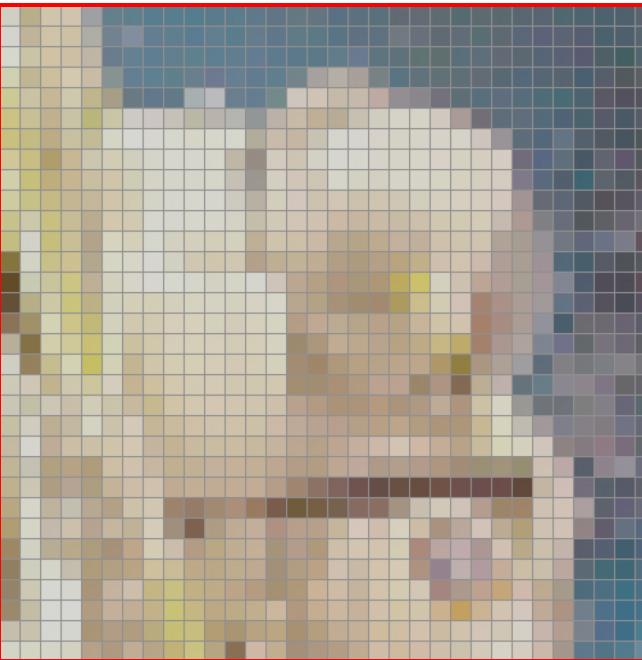
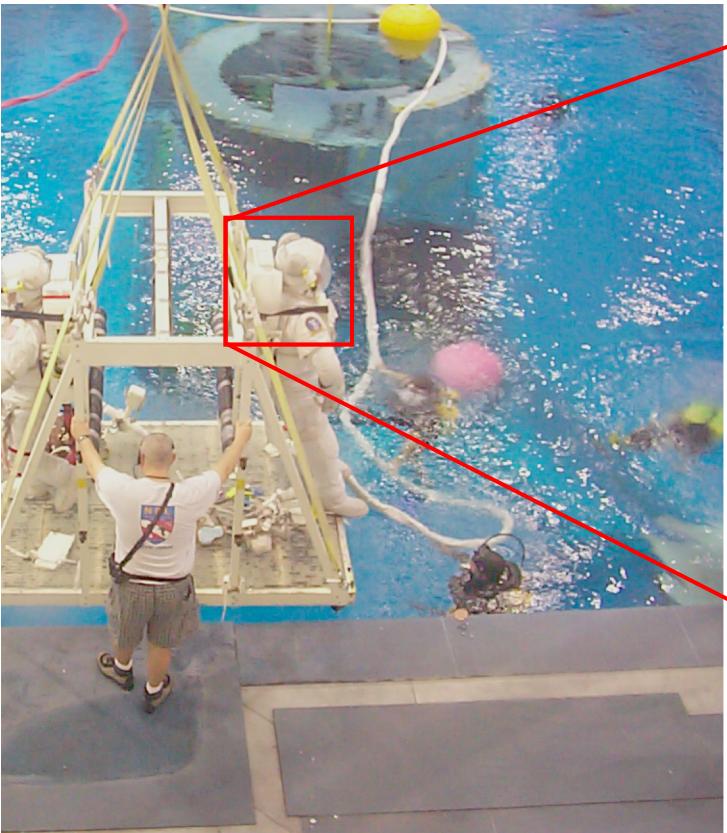
Pixelize the section to better see the effects.



apply a pixel grid

Moving-Window Transformations

Pixelize the section to better see the effects.

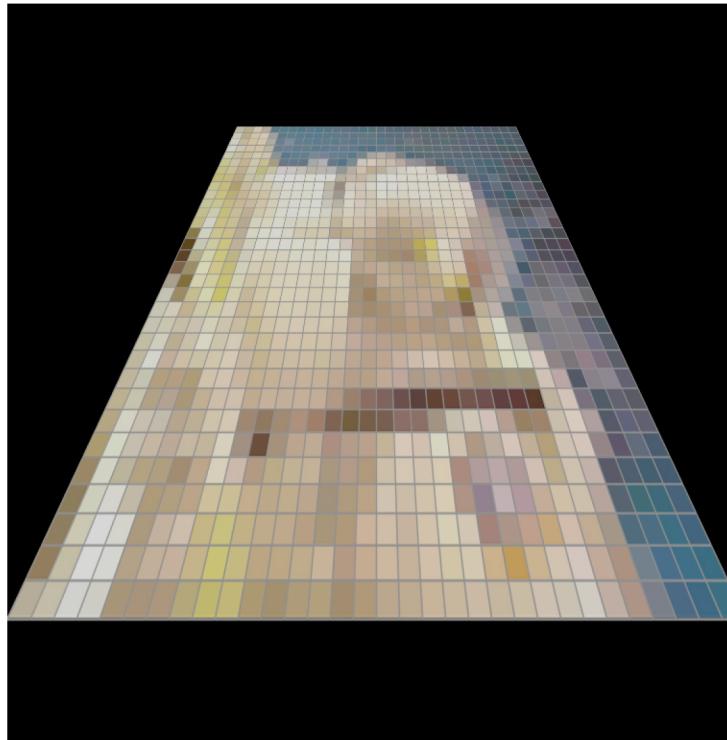


sample (average
in the squares).

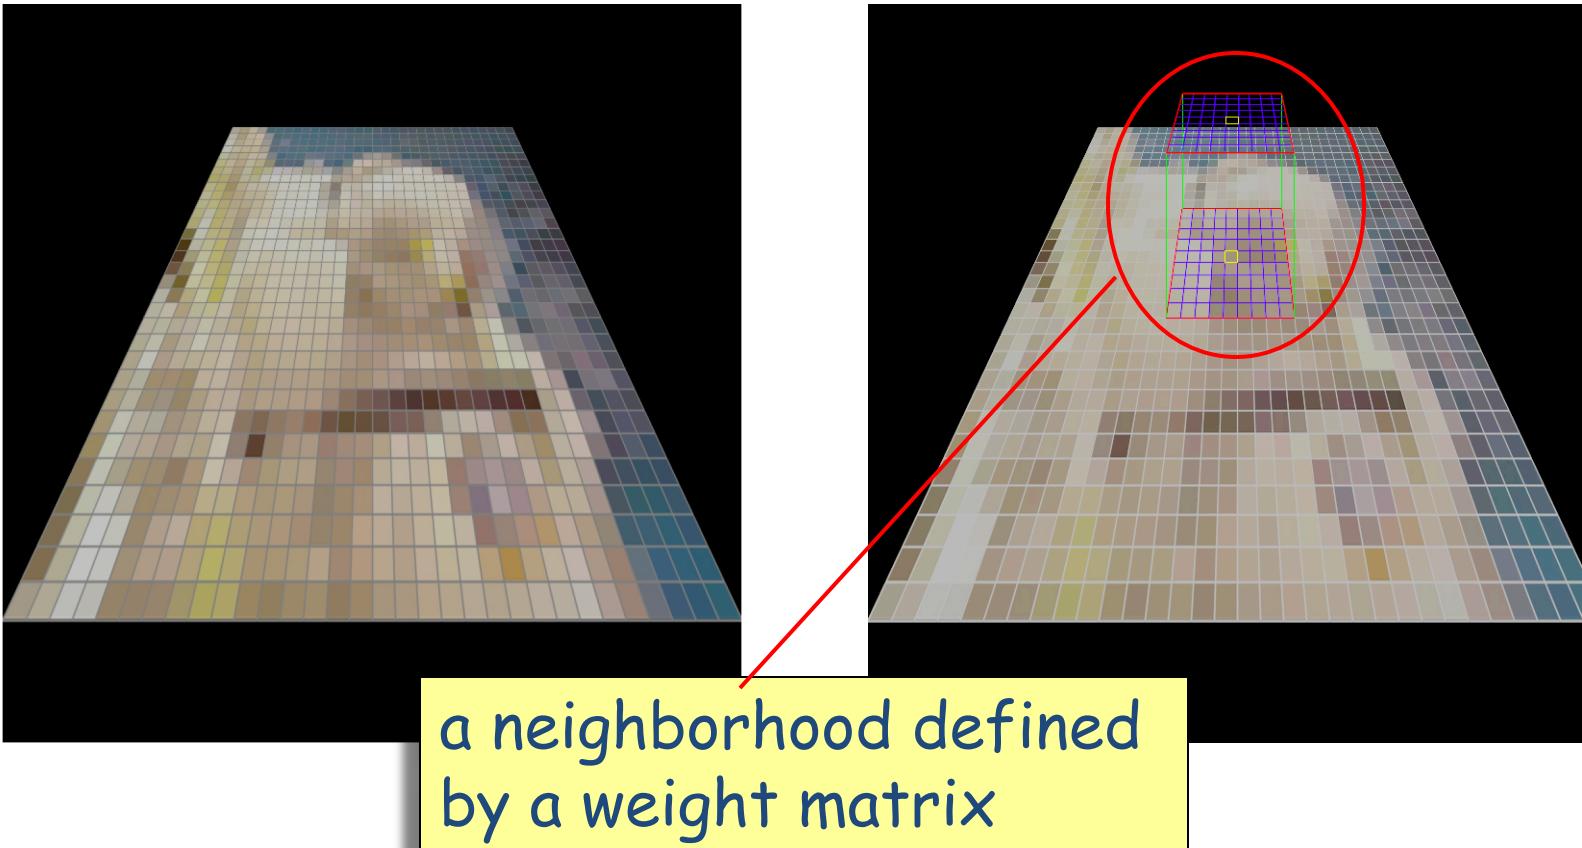
Moving-Window Transformations



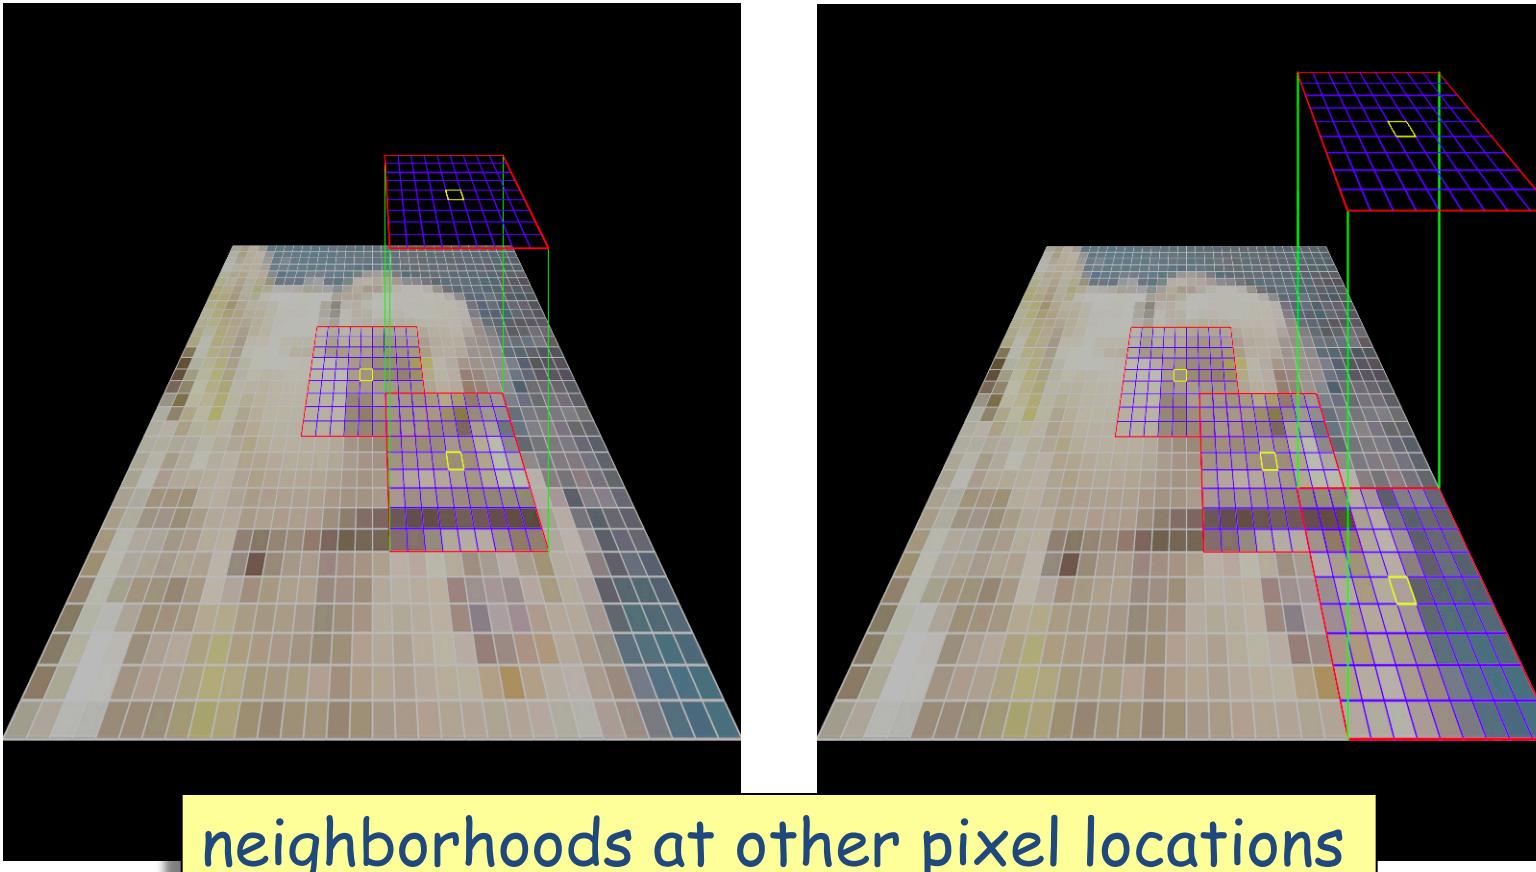
lets get some
perspective on
this



Moving-Window Transformations

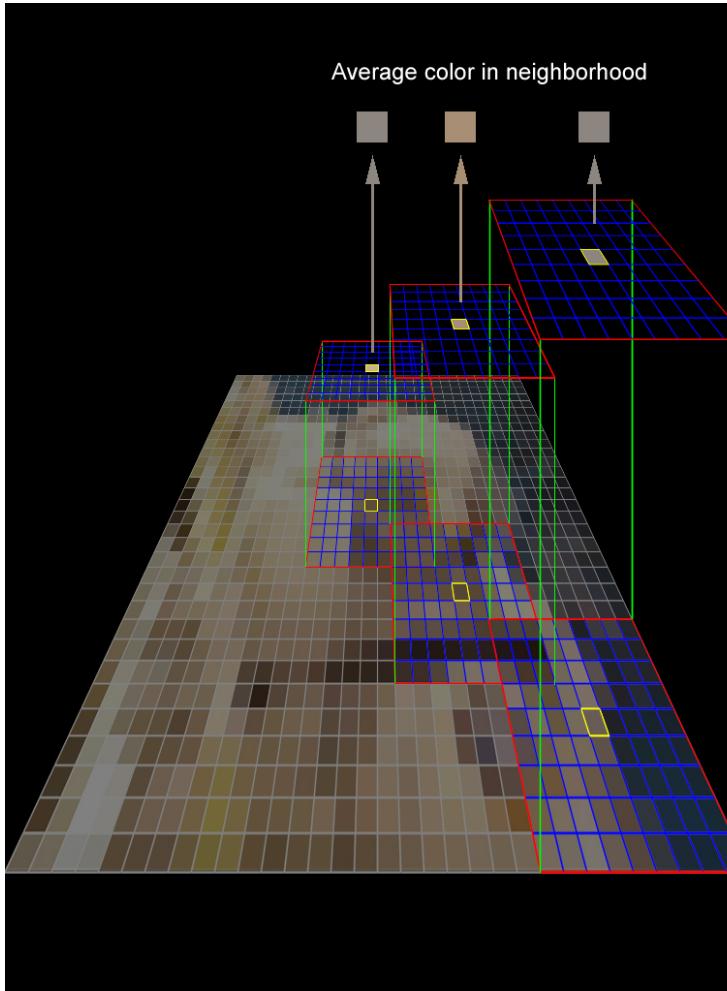


Moving-Window Transformations

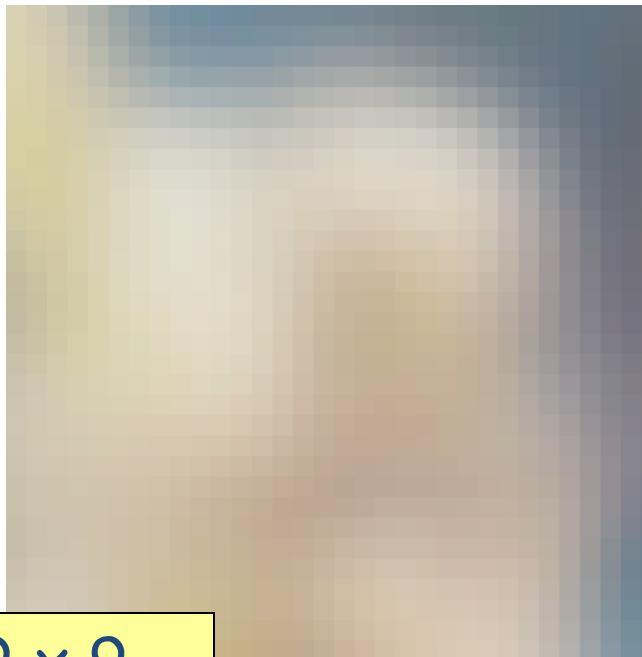
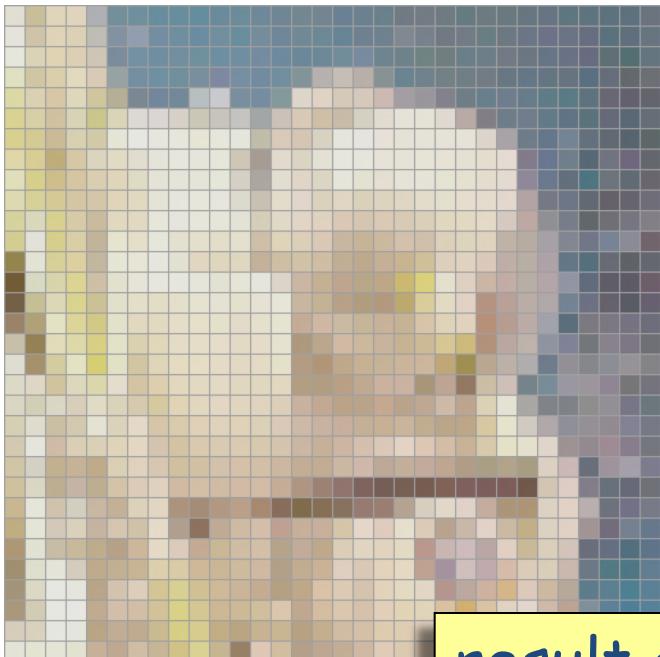


Linear Moving-Window Transformations (*i.e.* convolution)

The output of the transform at each pixel is the (weighted) average of the pixels in the neighborhood.



Moving-Window Transformations



result of a 9×9
uniform averaging

Convolution: Mathematical Representation

If a MW transformation is *linear* then it is a *convolution*:

$$\mathbf{J}(r, c) = [\mathbf{I} * \mathbf{h}](r, c) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} \mathbf{I}(r - \rho, c - \chi) \mathbf{h}(\rho, \chi) d\rho d\chi,$$

for a real image ($\mathbf{I}: R \times R \rightarrow R$), or for a digital image ($\mathbf{I}: Z \times Z \rightarrow Z$):

$$\mathbf{J}(r, c) = [\mathbf{I} * \mathbf{h}](r, c) = \sum_{\rho=-s}^s \sum_{\chi=-d}^d \mathbf{I}(r - \rho, c - \chi) \mathbf{h}(\rho, \chi)$$

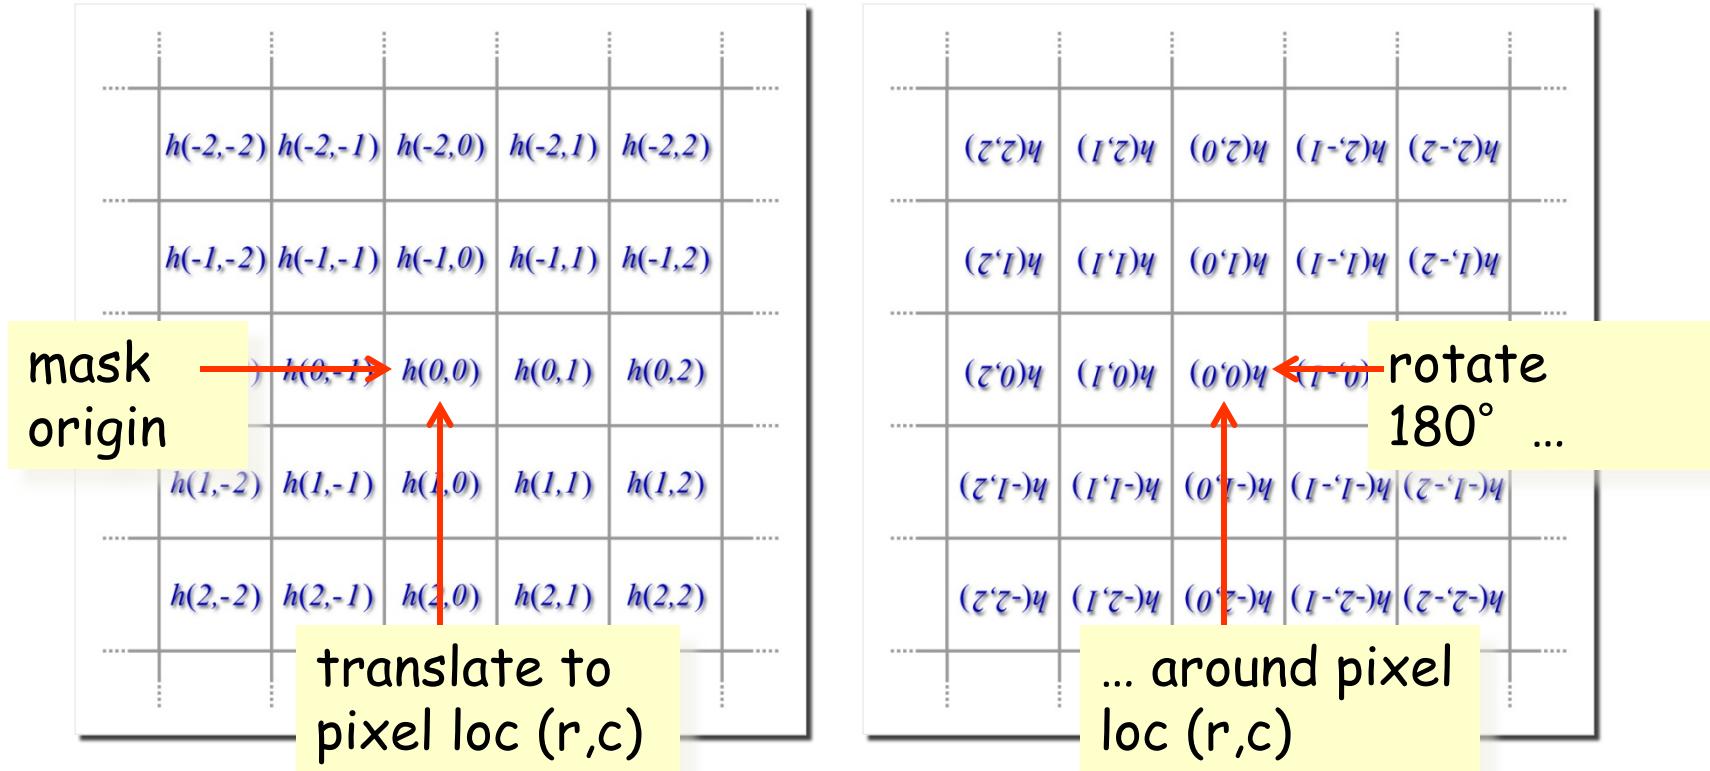


Convolution Mask (Weight Matrix)

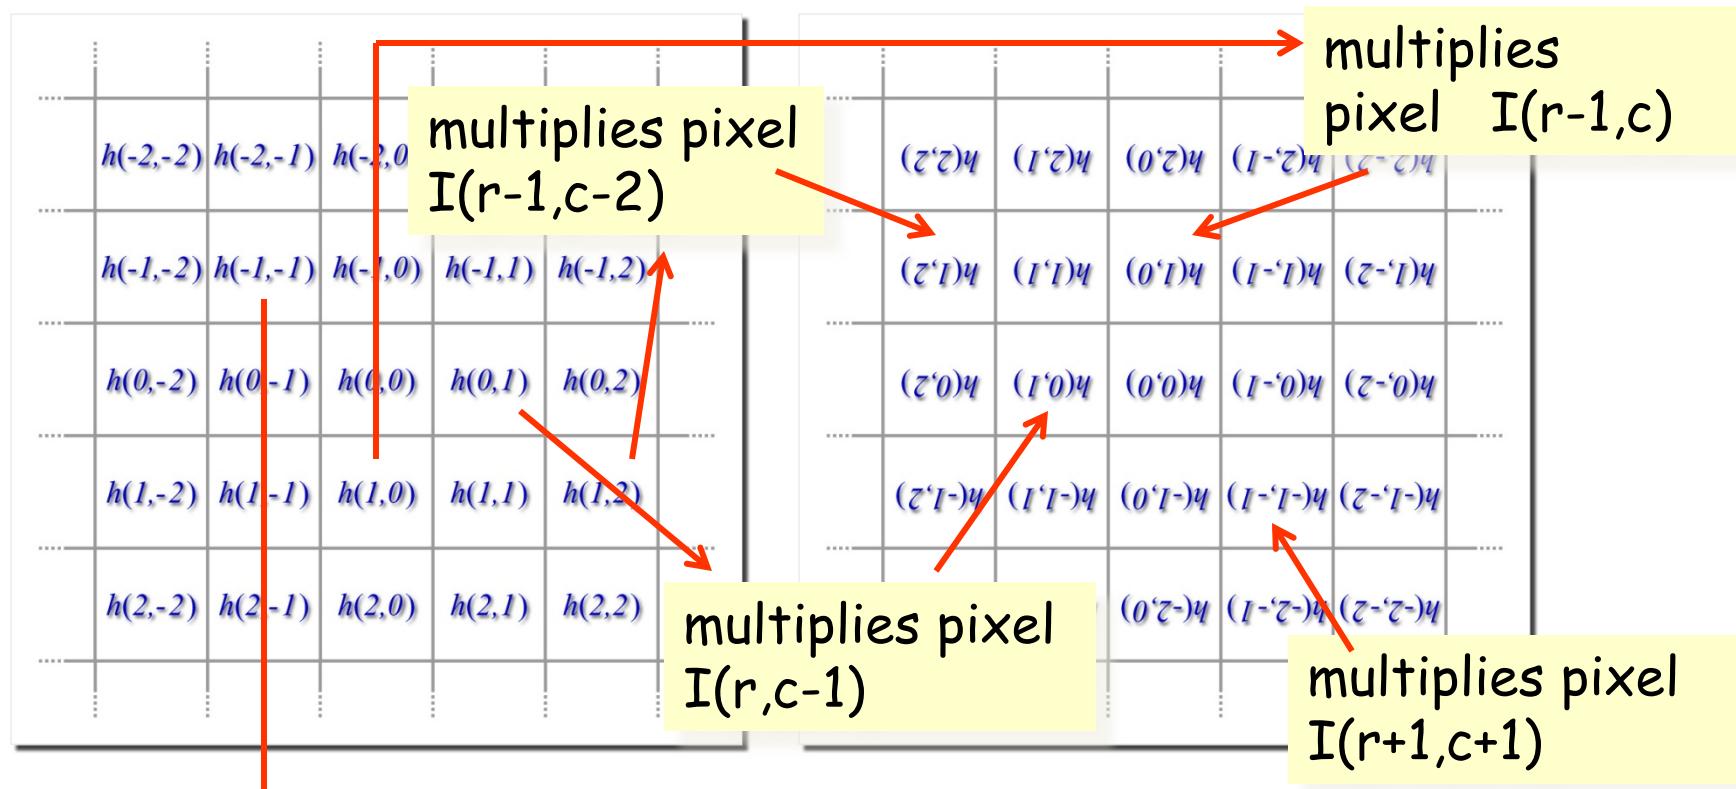
- The object, $h(\rho, \chi)$, in the equation is a weighting function, or in the discrete case, a rectangular matrix of numbers.
- The matrix is the moving window.
- Pixel (r,c) in the output image is the weighted sum of pixels from the original image in the neighborhood of (r,c) traced by the matrix.
- Each pixel in the neighborhood of (r,c) is multiplied by the corresponding matrix value — after the matrix is rotated by 180° .
- The sum of those products is the value of pixel (r,c) in the output image



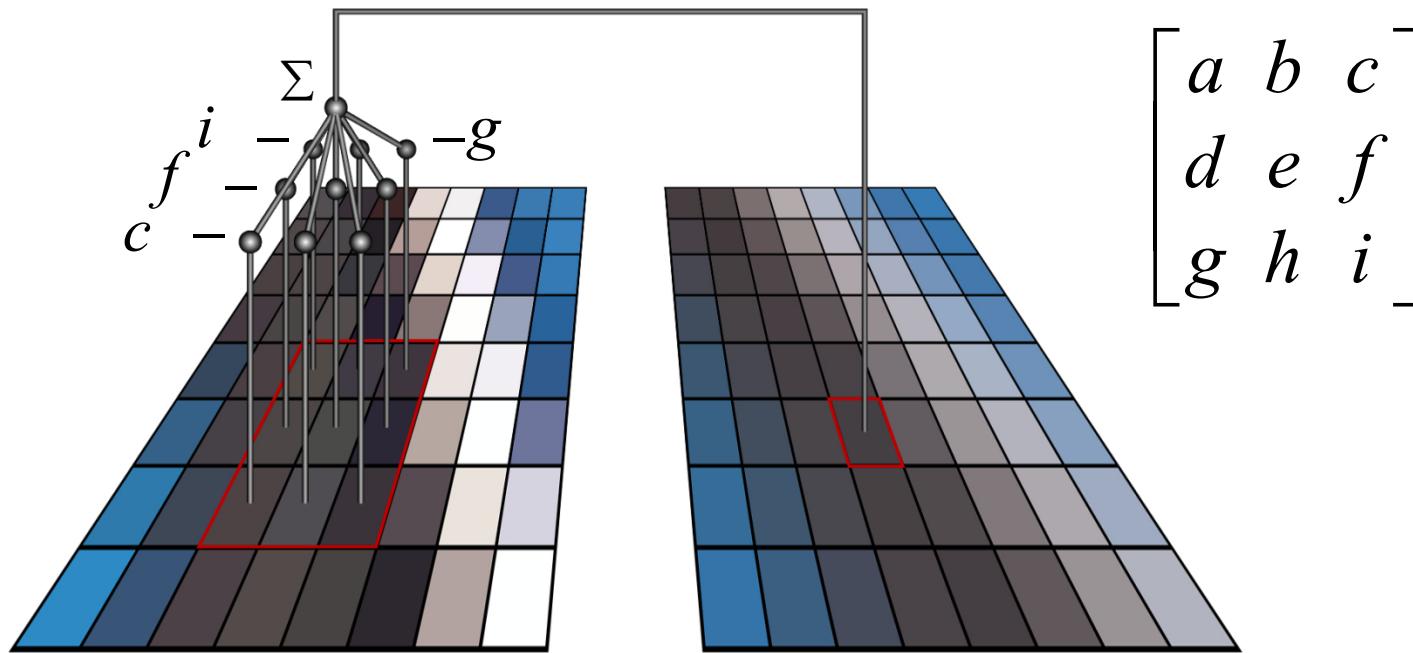
Convolution Masks: Moving Window



Convolution Masks: Moving Window

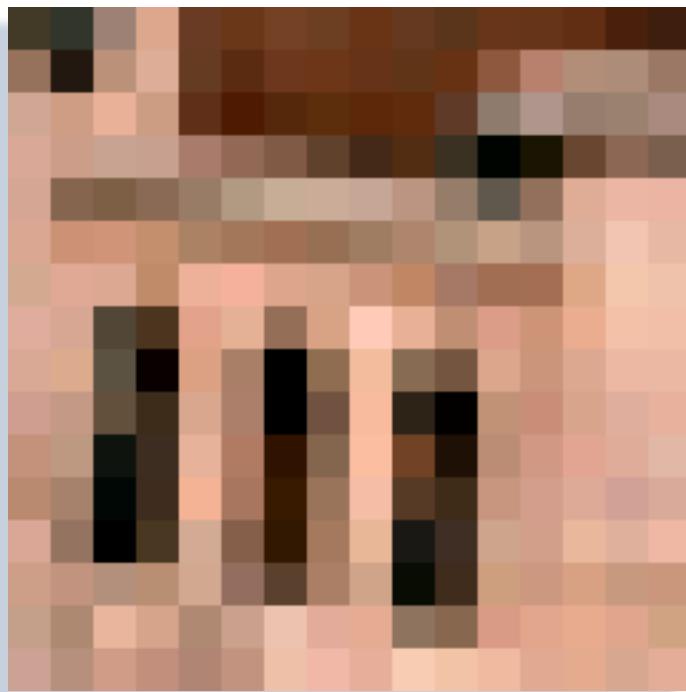


Convolution by Moving Window

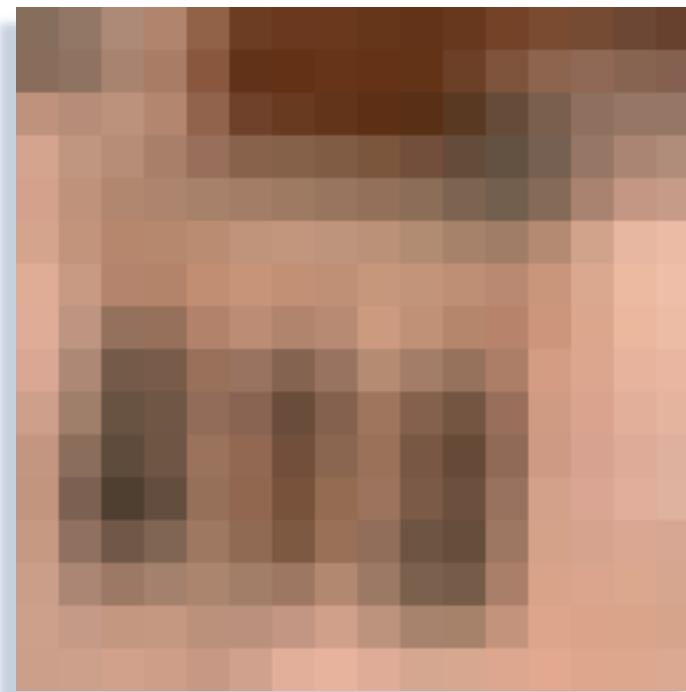


Moving Window Transform: Example

Another example

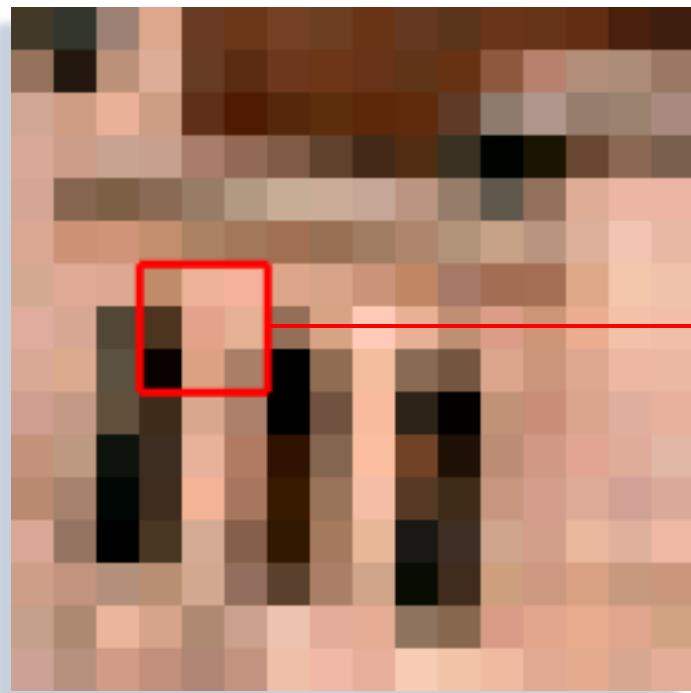


original

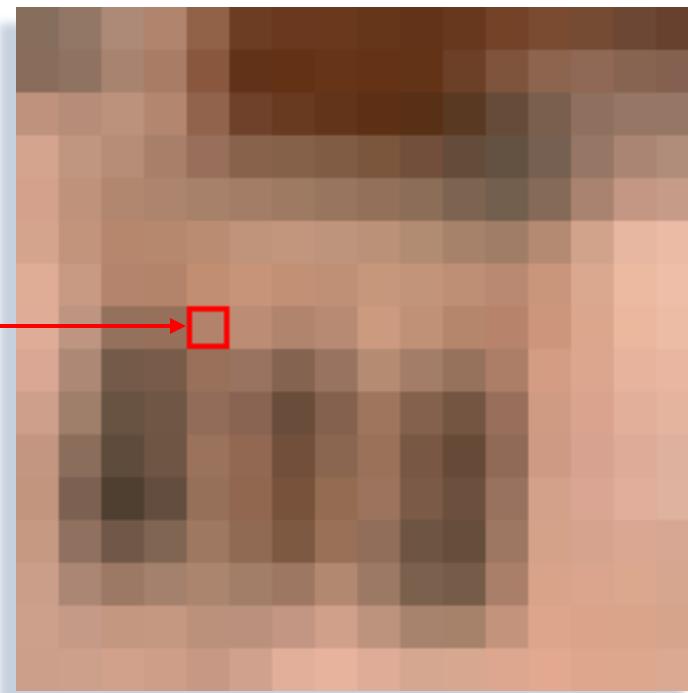


3x3 average

Moving Window Transform: Example

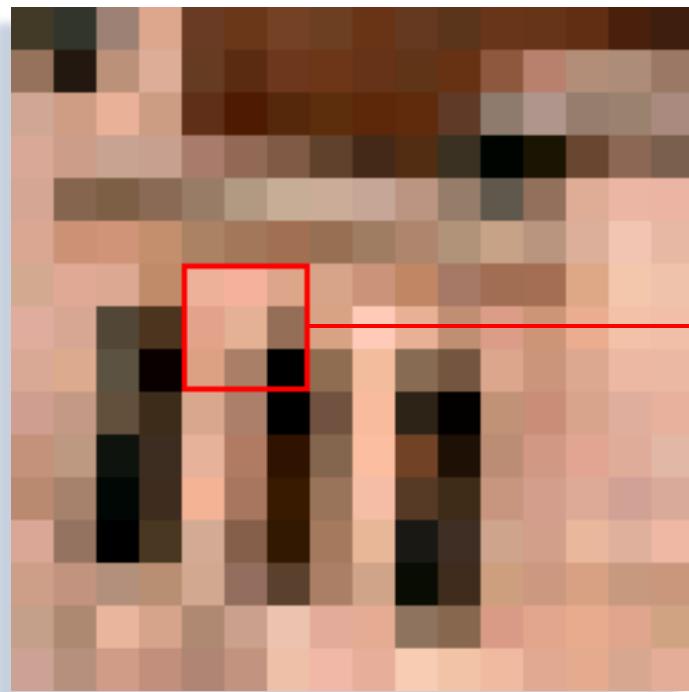


original

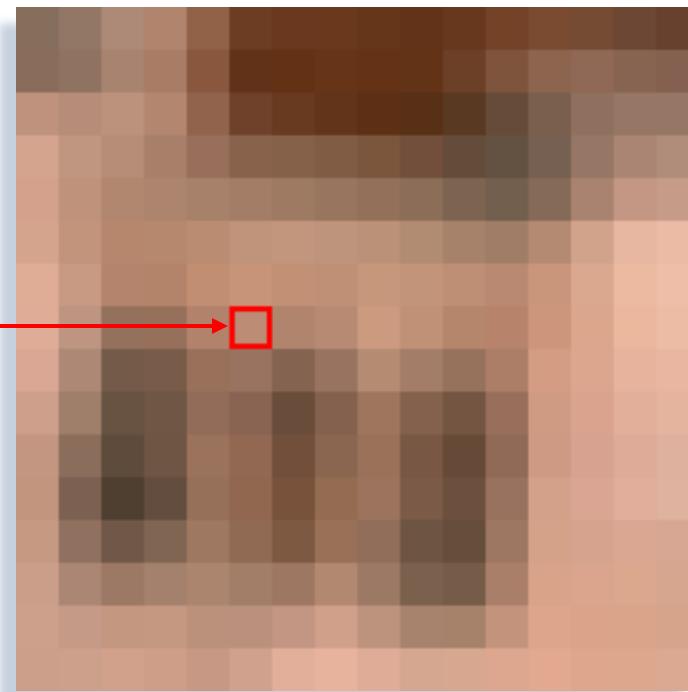


3x3 average

Moving Window Transform: Example

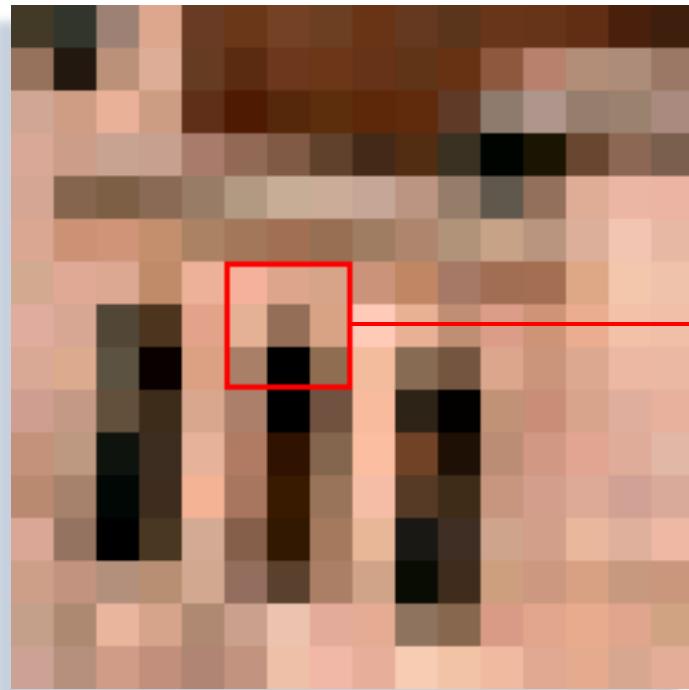


original

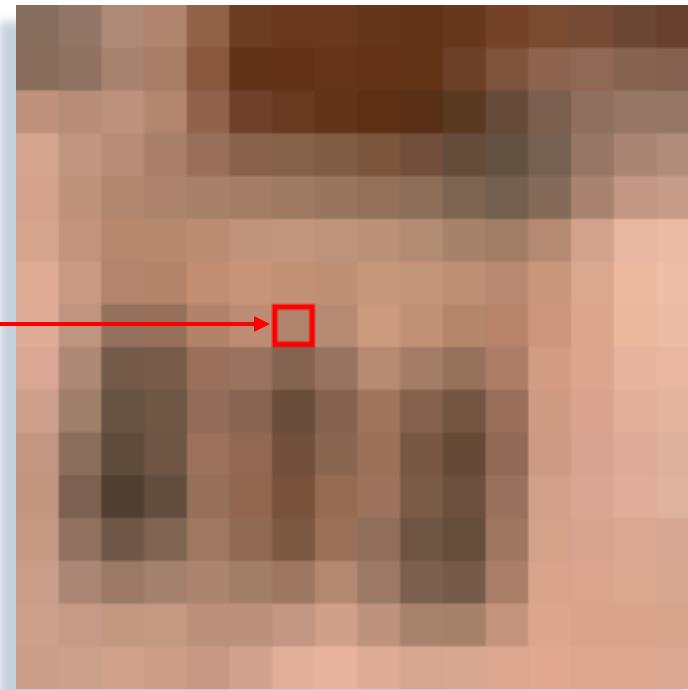


3x3 average

Moving Window Transform: Example

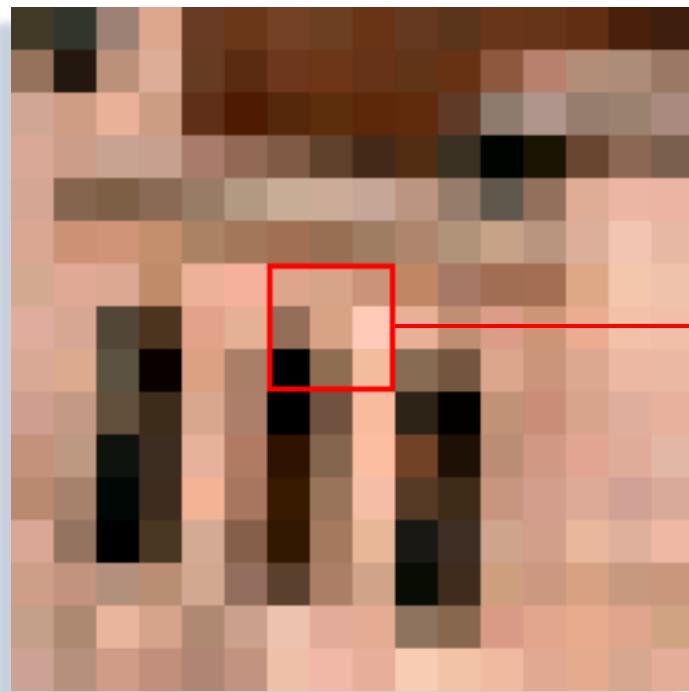


original

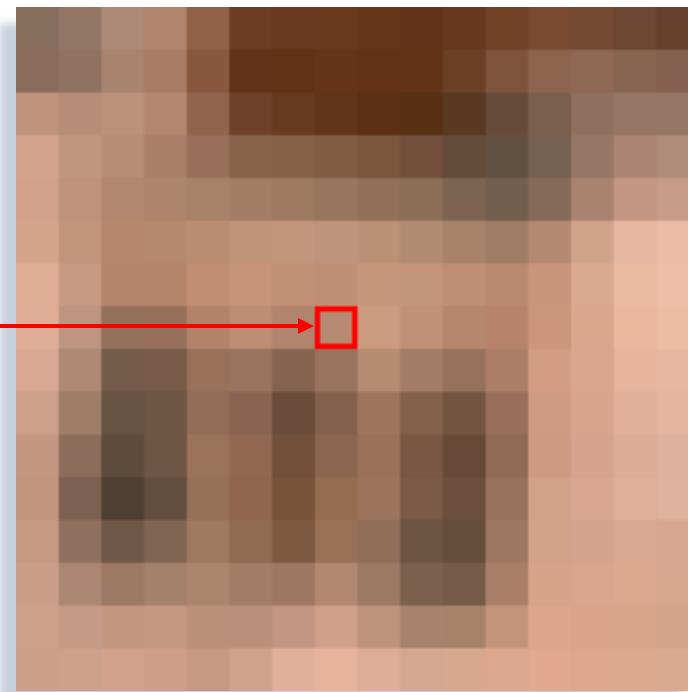


3x3 average

Moving Window Transform: Example

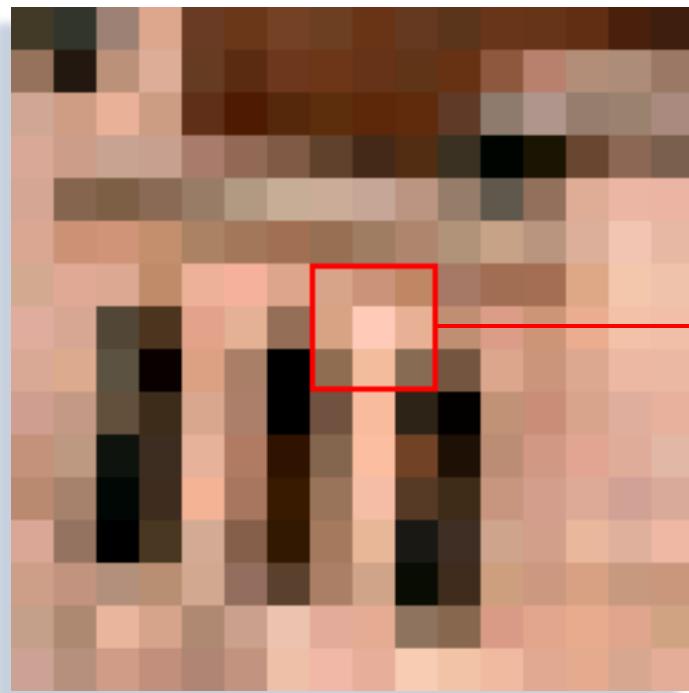


original

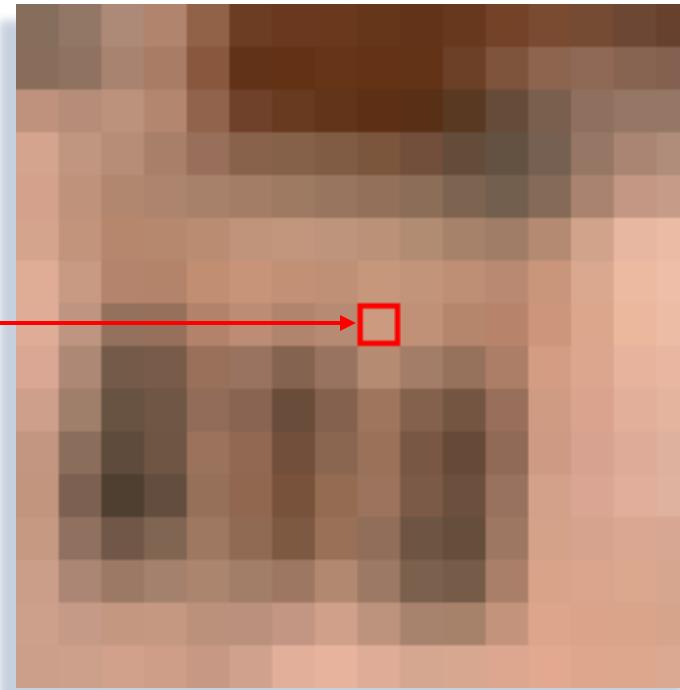


3x3 average

Moving Window Transform: Example

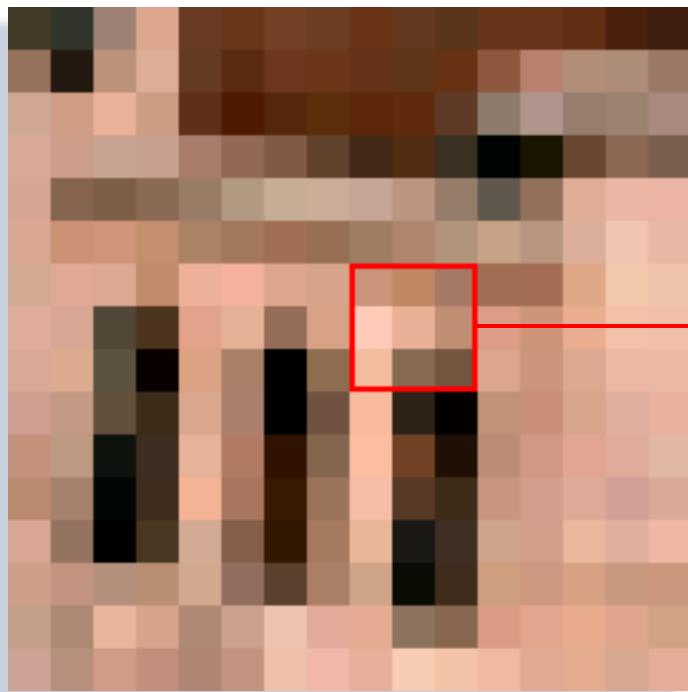


original

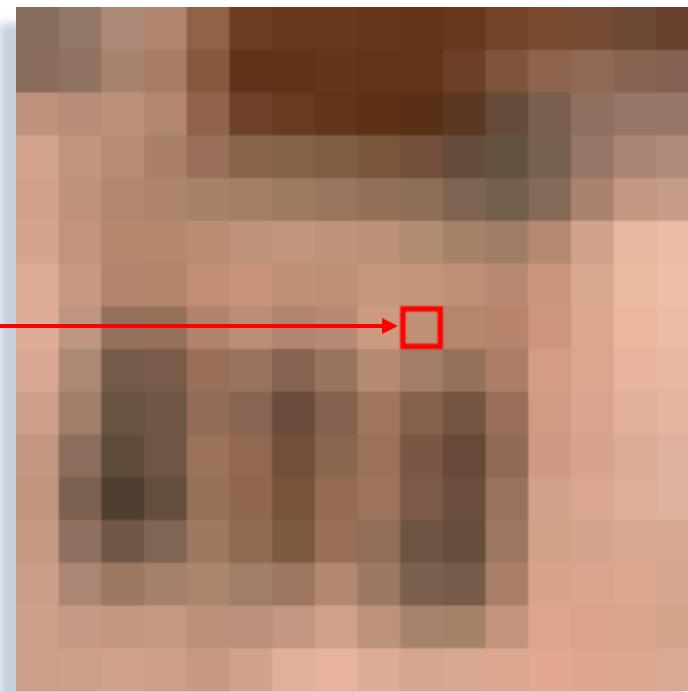


3×3 average

Moving Window Transform: Example

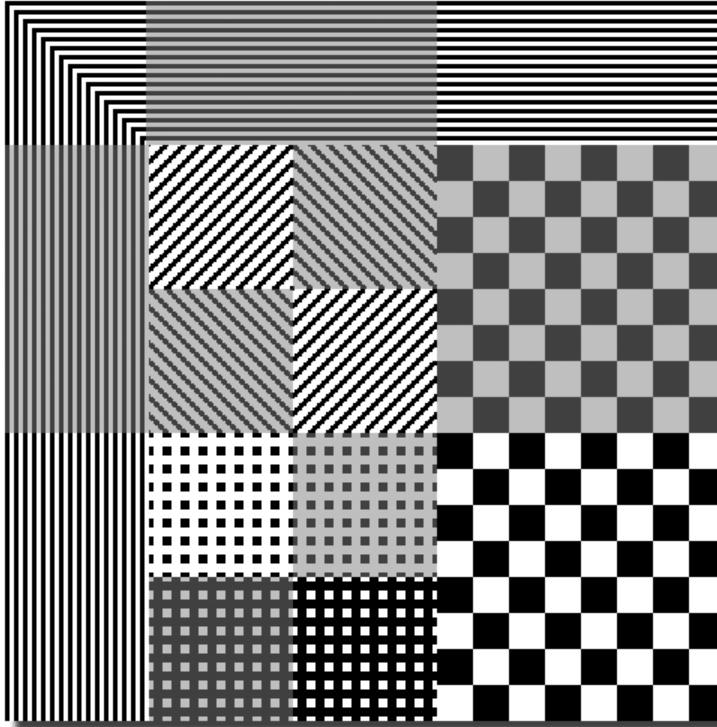


original



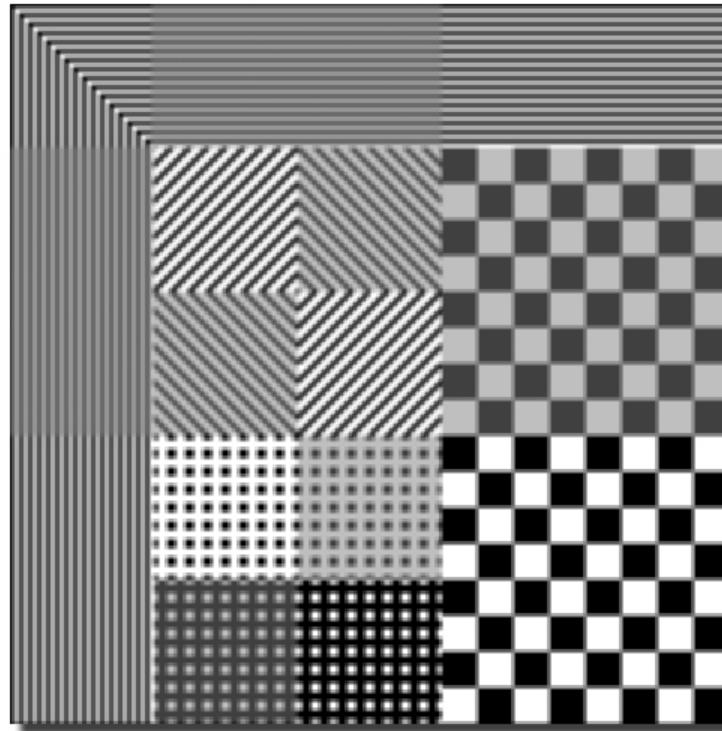
3x3 average

Convolution Examples: Original Images



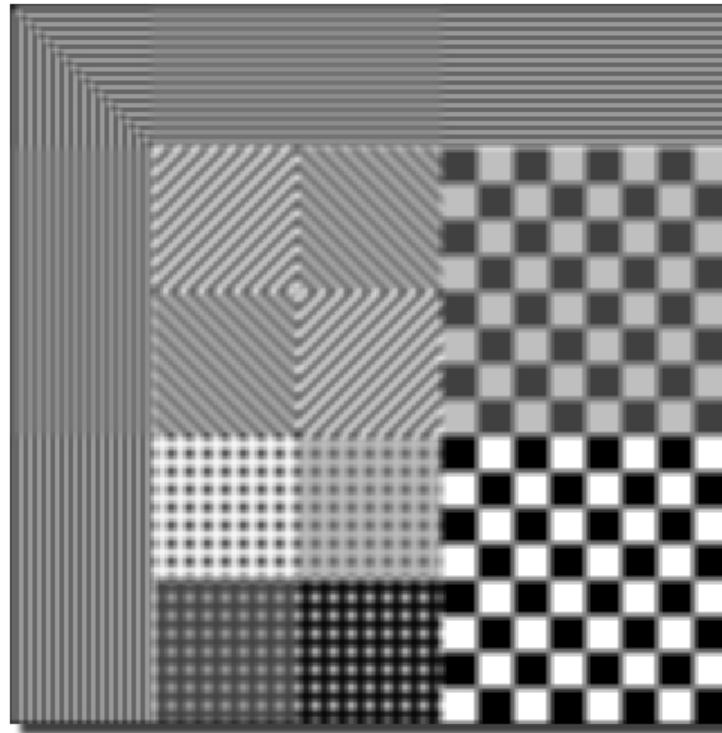
Convolution Examples: 3×3 Blur

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$



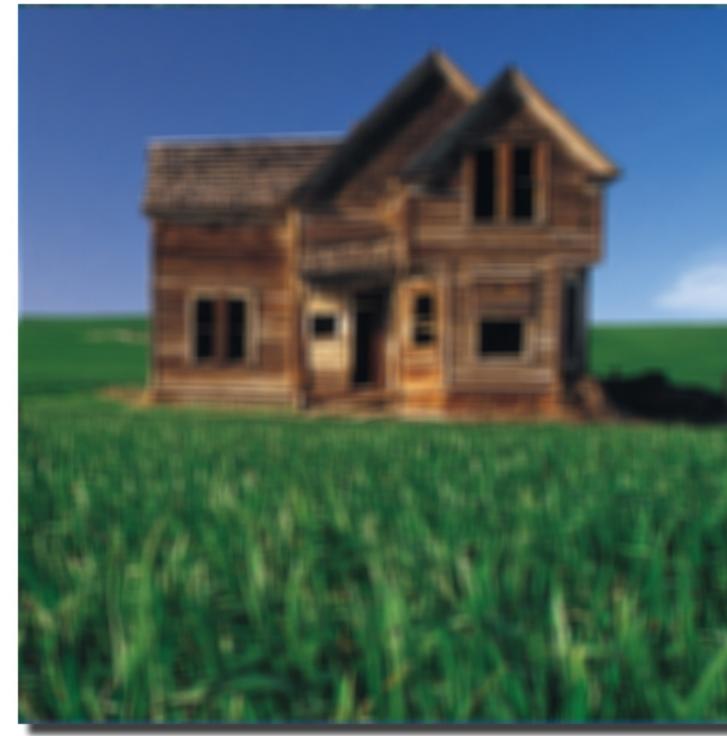
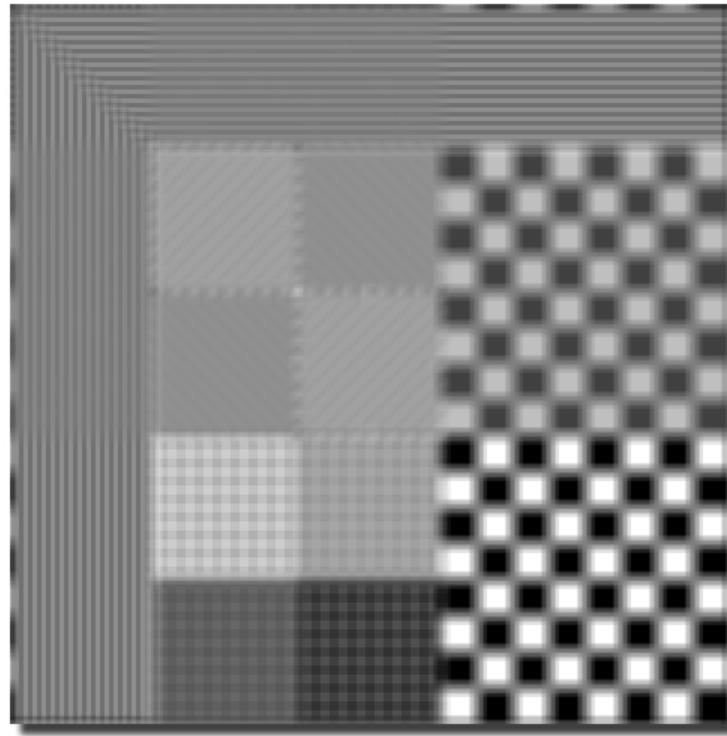
Convolution Examples: 5×5 Blur

$$\frac{1}{25} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$

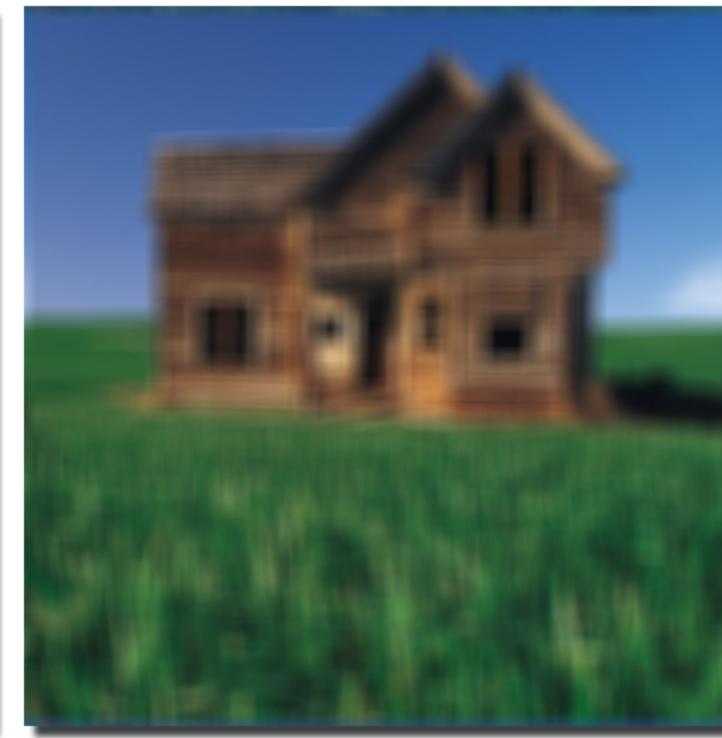
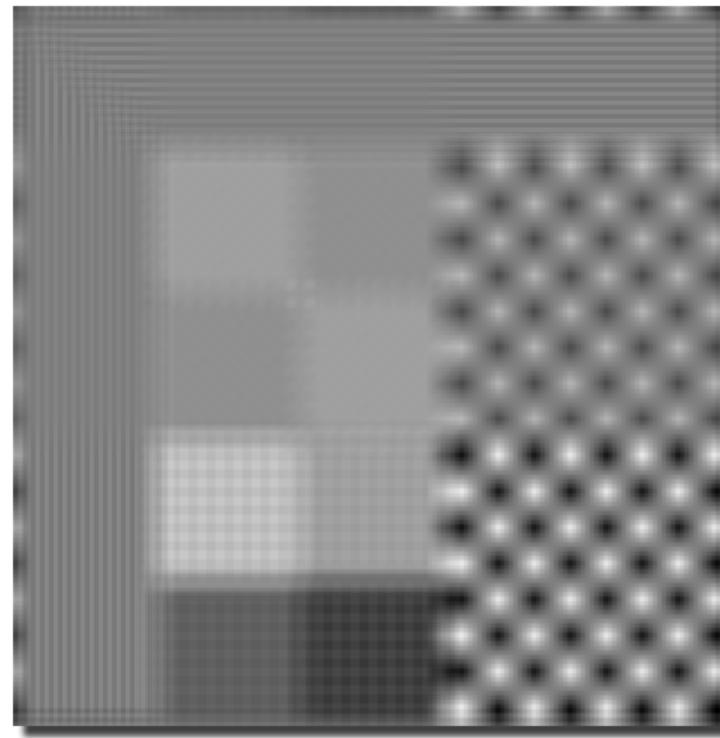


Convolution Examples: 9×9 Blur

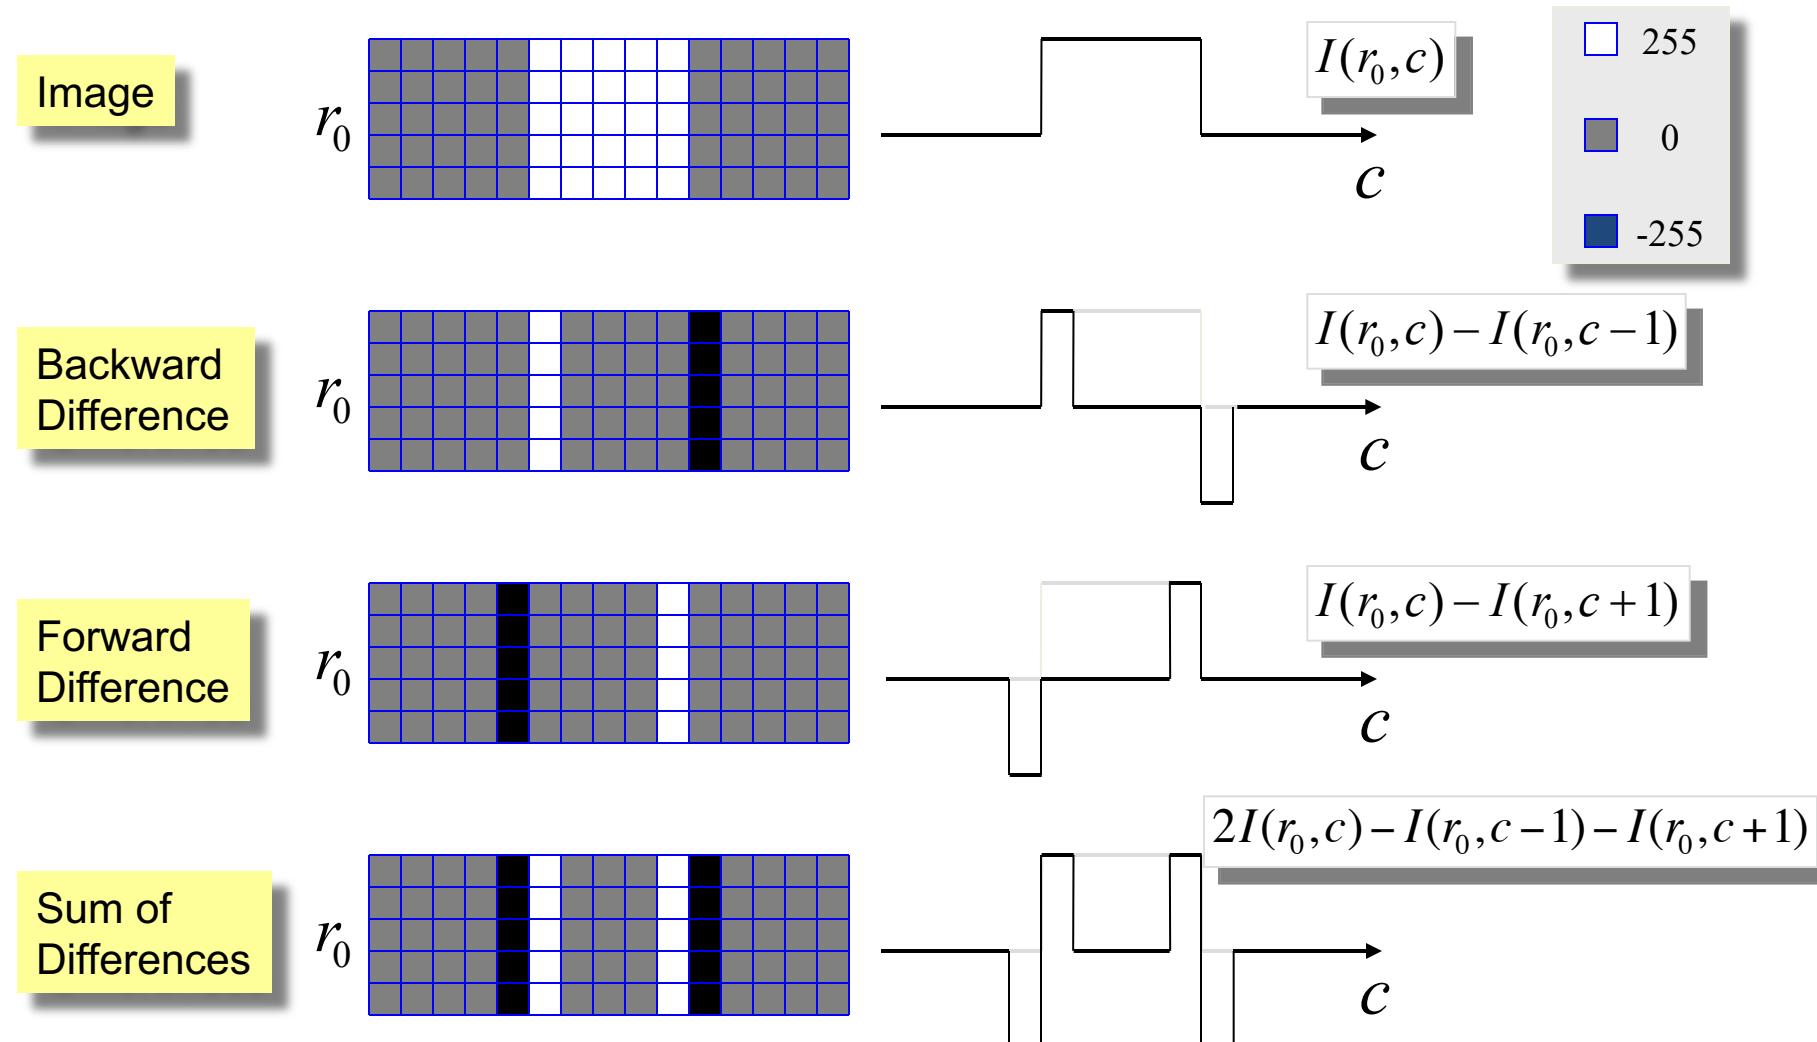
$$\frac{1}{81} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 & 1 \end{bmatrix}$$



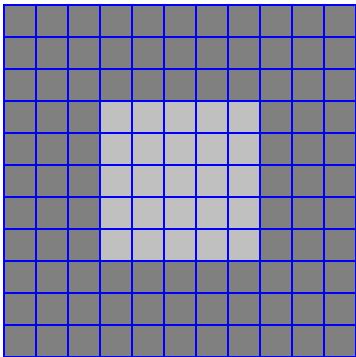
Convolution Examples: 17×17 Blur



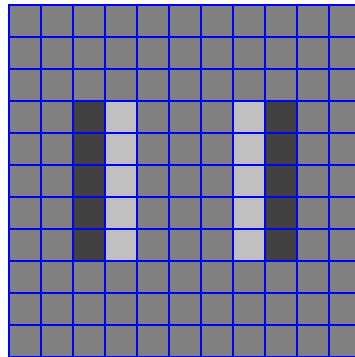
Vertical Edge Detection



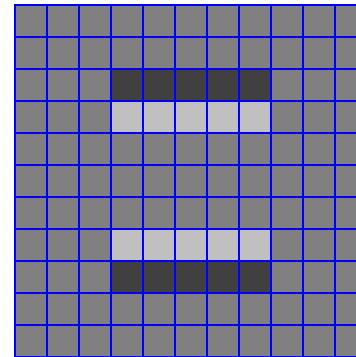
Symmetric Edge Detection



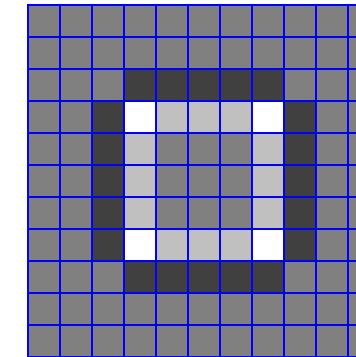
$I(r,c)$



$$2I(r,c) - I(r,c-1) - I(r,c+1)$$



$$2I(r,c) - I(r-1,c) - I(r+1,c)$$



$$4I(r,c) - I(r-1,c) - I(r+1,c) - I(r,c-1) - I(r,c+1)$$

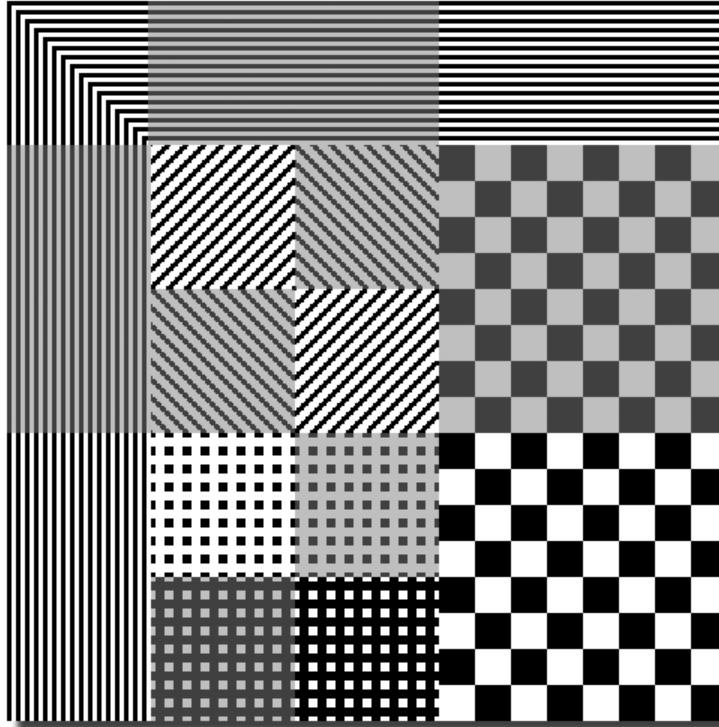
510
255
0
-255

-1	2	-1

	-1	
	2	
	-1	

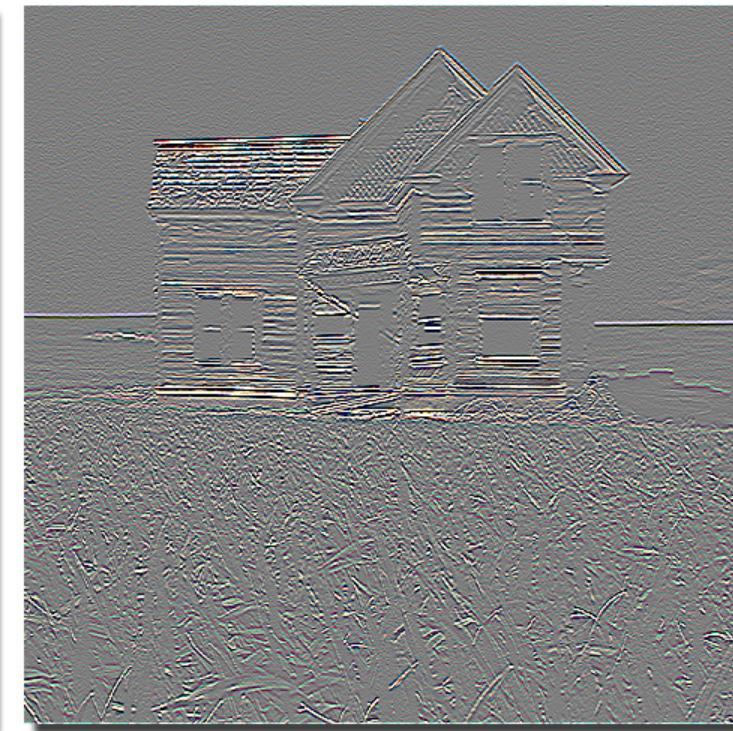
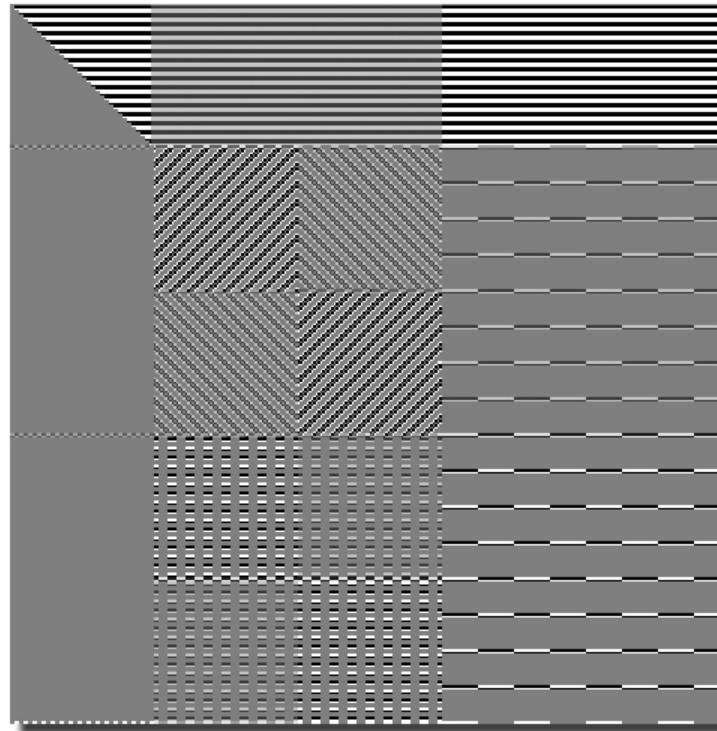
	-1	
-1	4	-1
	-1	

Convolution Examples: Original Images



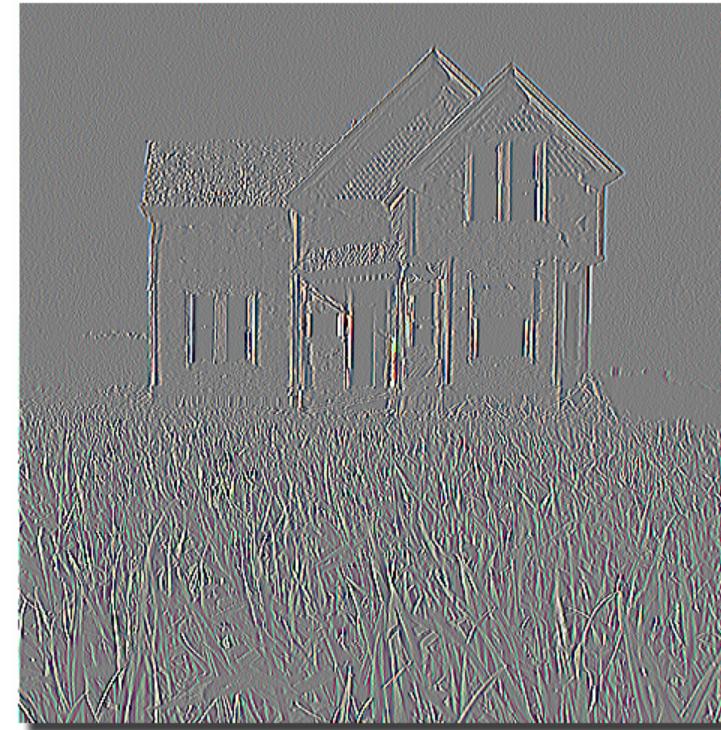
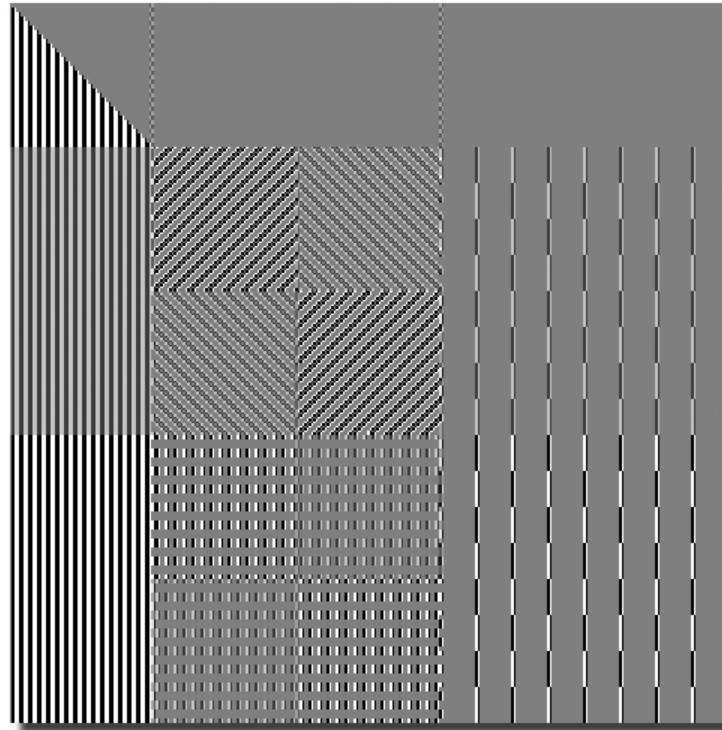
Convolution Examples: Vertical Difference

$$\begin{bmatrix} -1 \\ 2 \\ -1 \end{bmatrix}$$



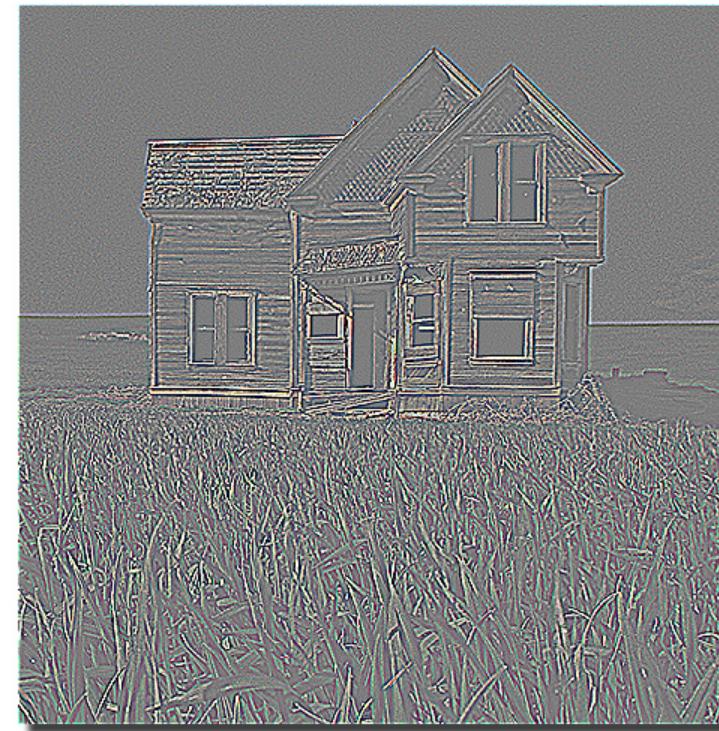
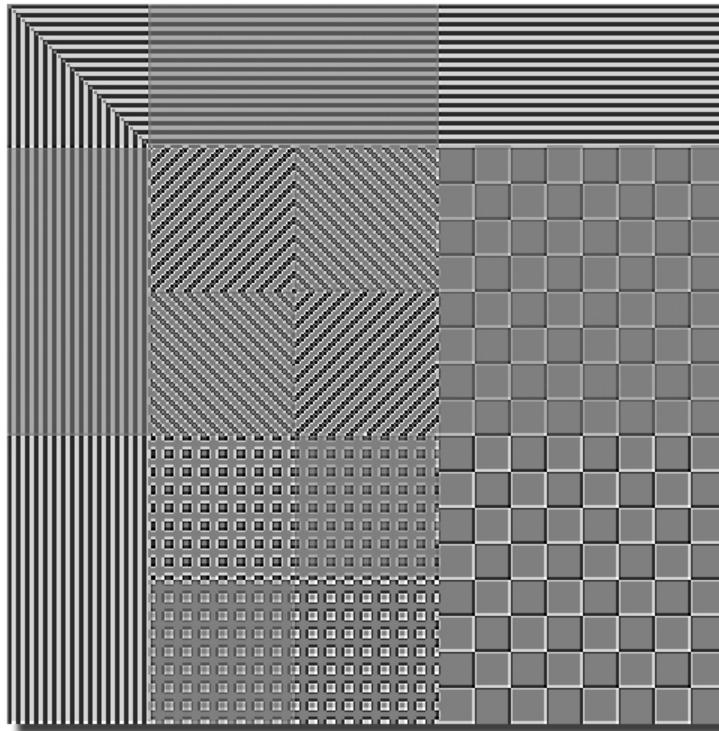
Convolution Examples: Horizontal Difference

$$[-1 \quad 2 \quad -1]$$



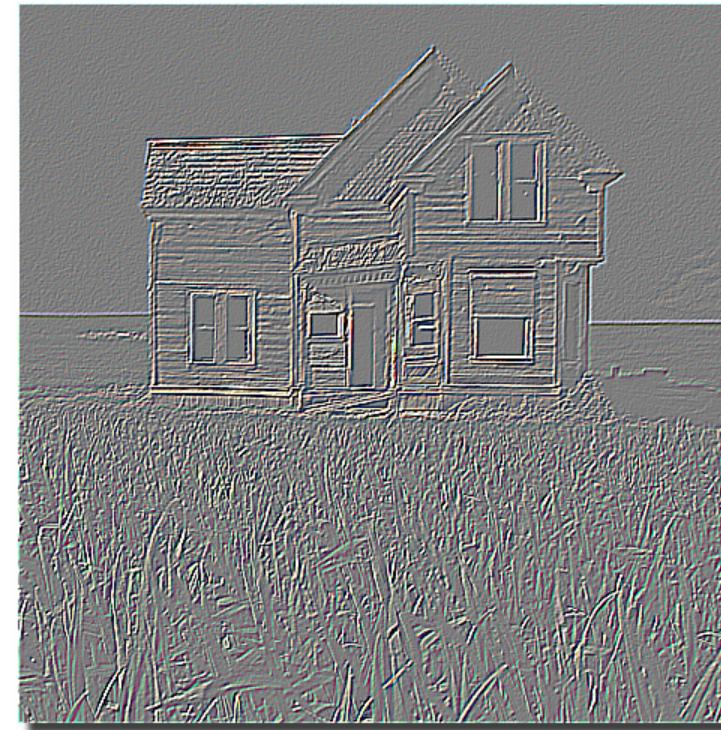
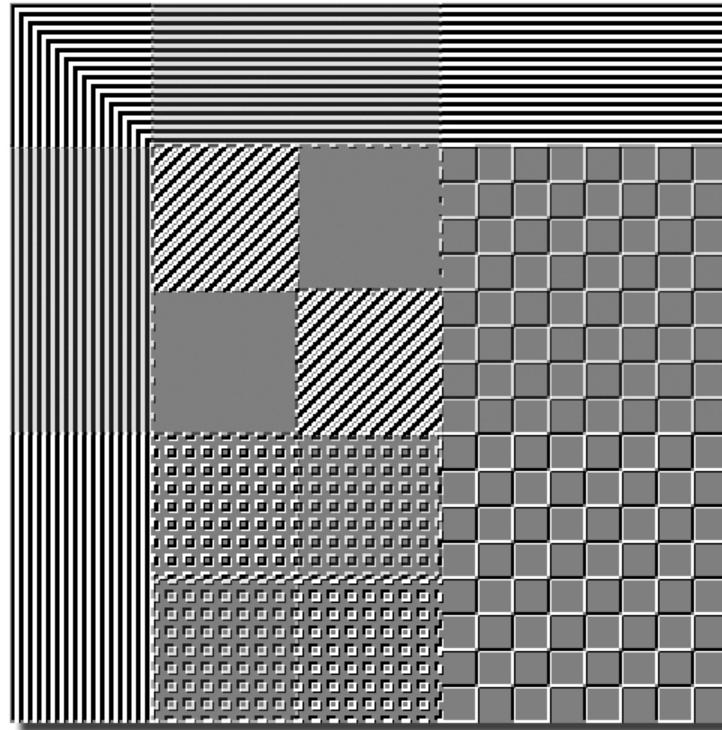
Convolution Examples: H + V Diff.

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



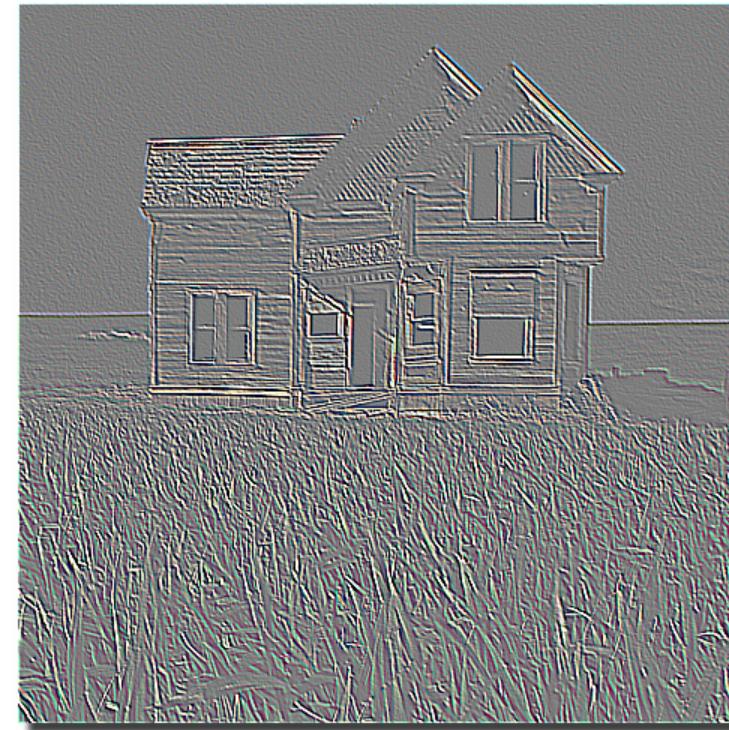
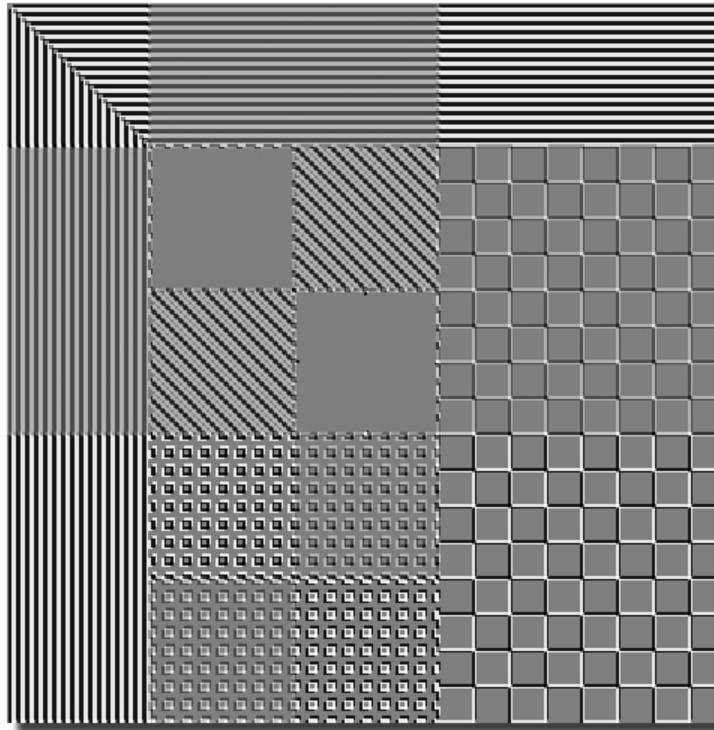
Convolution Examples: Diagonal Difference

$$\begin{bmatrix} -1 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & -1 \end{bmatrix}$$



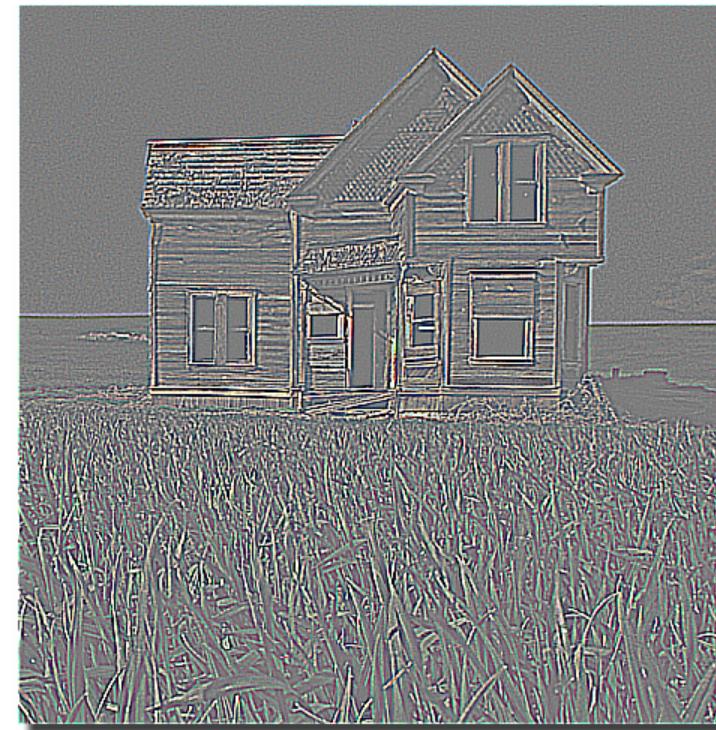
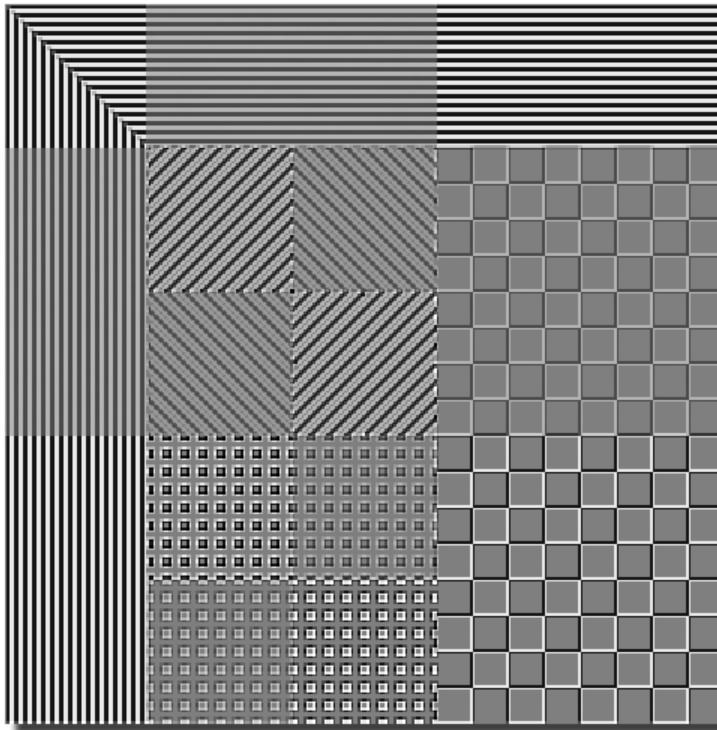
Convolution Examples: Diagonal Difference

$$\begin{bmatrix} 0 & 0 & -1 \\ 0 & 2 & 0 \\ -1 & 0 & 0 \end{bmatrix}$$



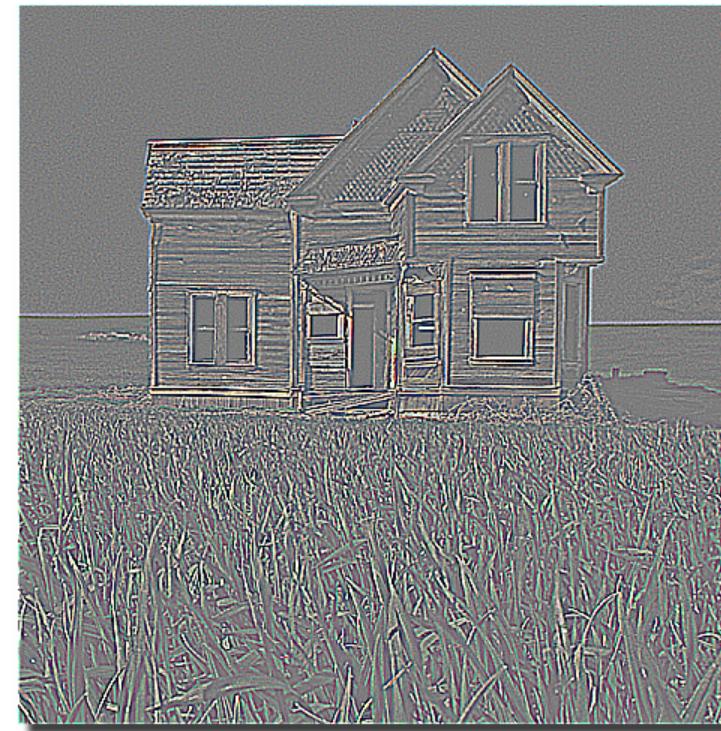
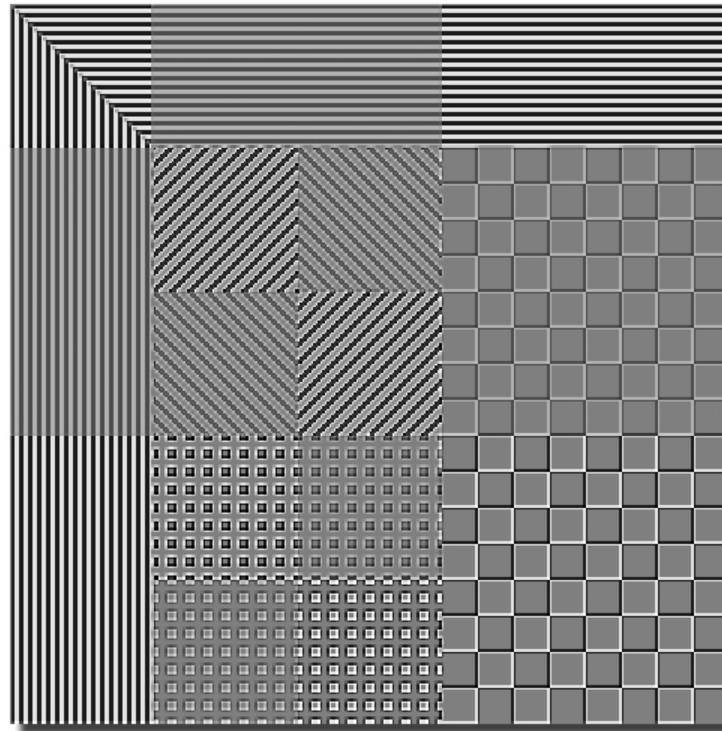
Convolution Examples: D + D Difference

$$\begin{bmatrix} -1 & 0 & -1 \\ 0 & 4 & 0 \\ -1 & 0 & -1 \end{bmatrix}$$



Convolution Examples: H + V + D Diff.

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$



The Median Filter

- Returns the median value of the pixels in a neighborhood
- Is non-linear
- Is similar to a uniform blurring filter which returns the mean value of the pixels in a neighborhood of a pixel
- Unlike a mean value filter the median tends to preserve step edges



Median Filter: General Definition

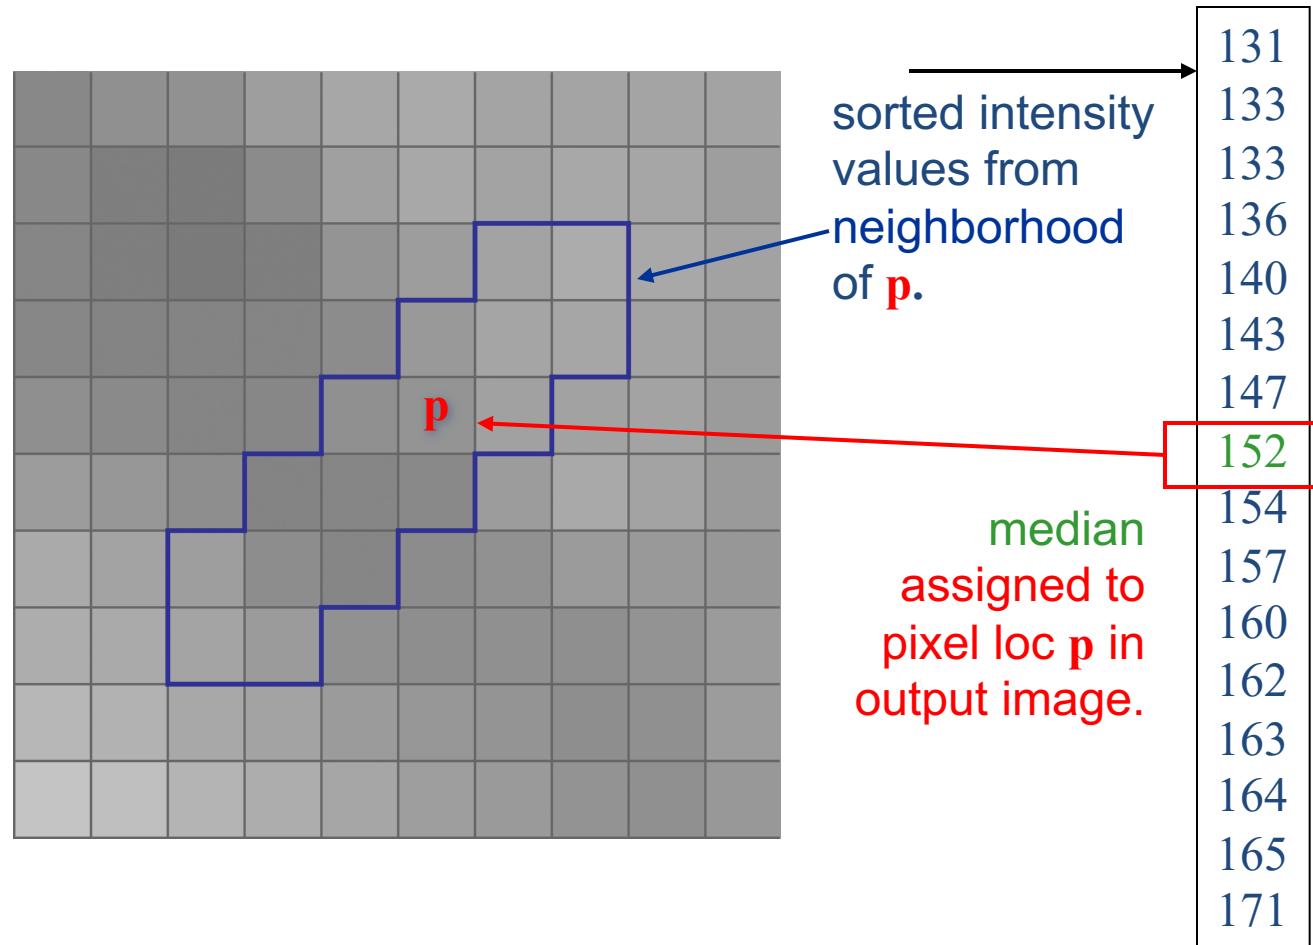
$$\text{med}\{I, Z\}(p) = \underset{q \in \text{supp}(Z+p)}{\text{median}} \{I(q)\}$$

This can be computed as follows:

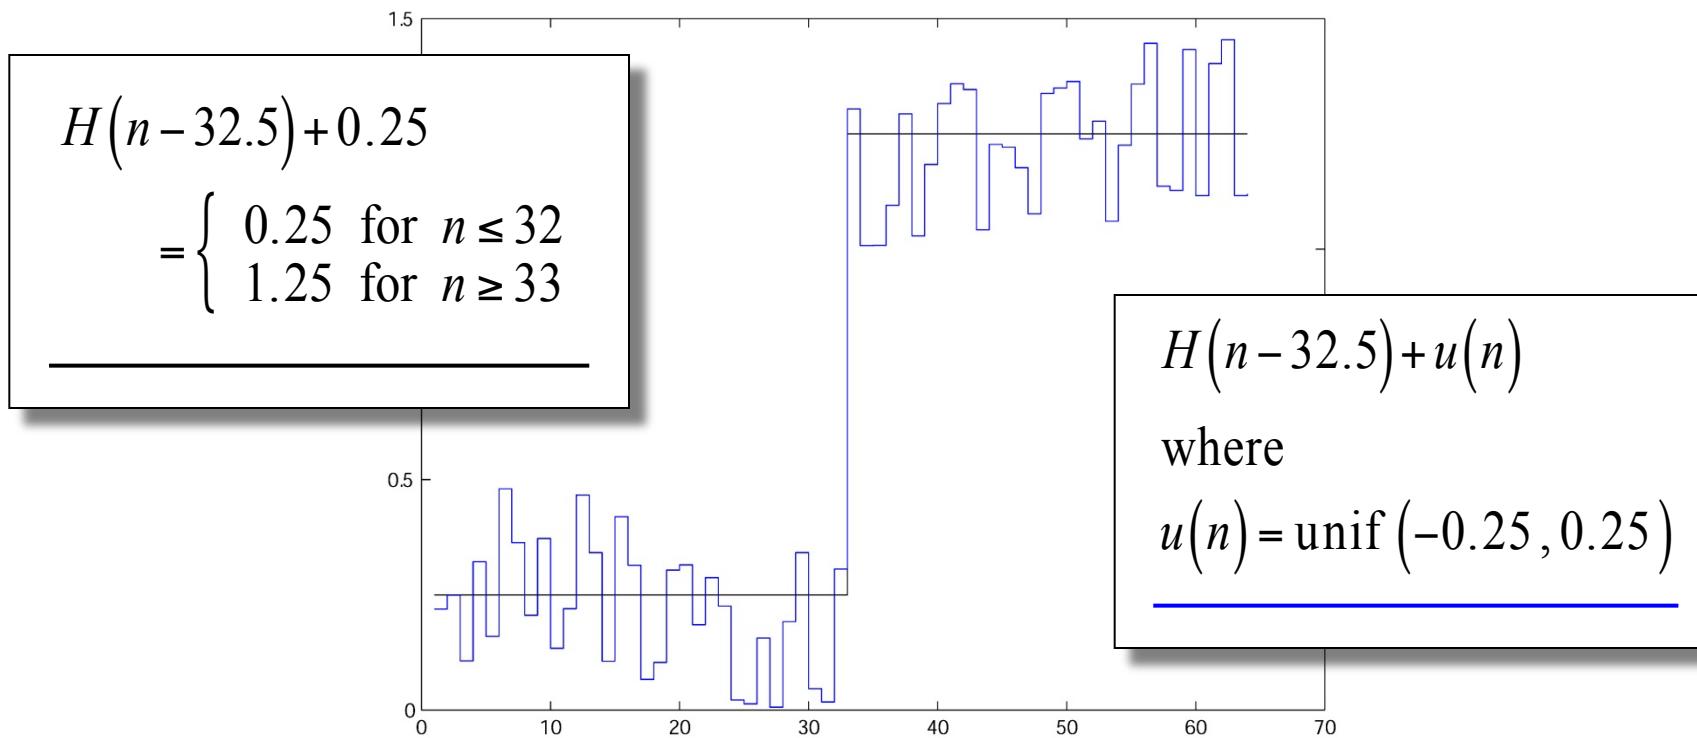
1. Let I be a monochrome (1-band) image.
2. Let Z define a neighborhood of arbitrary shape.
3. At each pixel location, $p = (r,c)$, in I ...
4. ... select the n pixels in the Z -neighborhood of p ,
5. ... sort the n pixels in the neighborhood of p , by
value, into a list $L(j)$ for $j = 1, \dots, n$.
6. The output value at p is $L(m)$, where $m = \lfloor \frac{n}{2} \rfloor + 1$



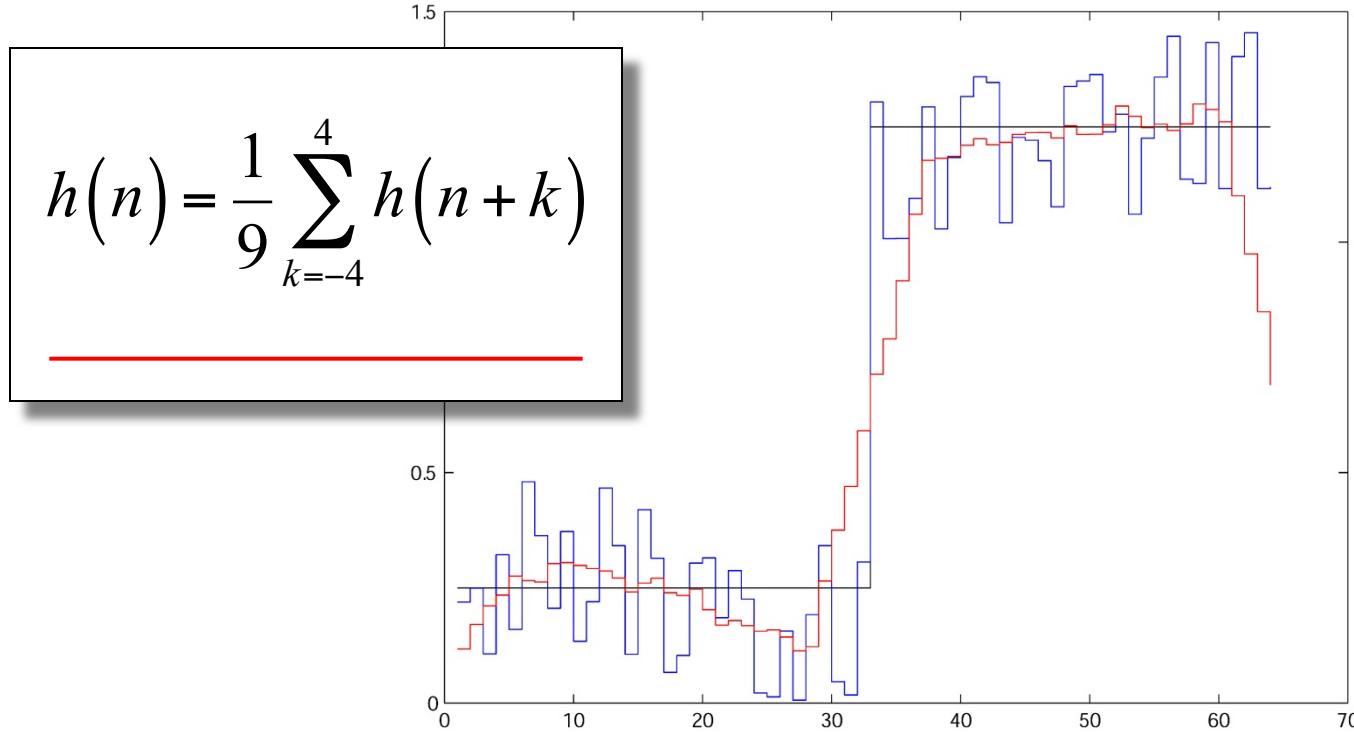
Median Filter: General Definition



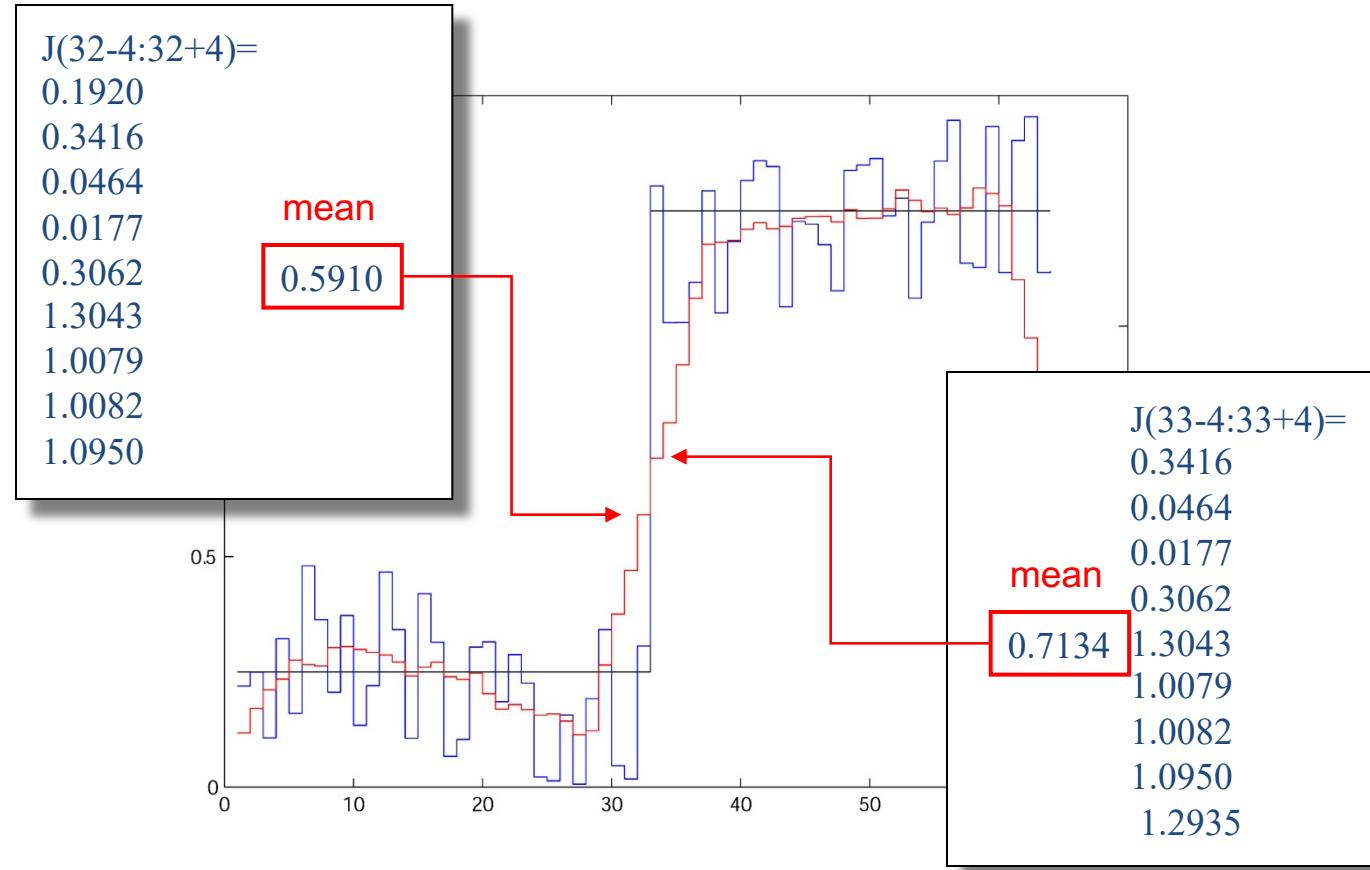
A Noisy Step Edge



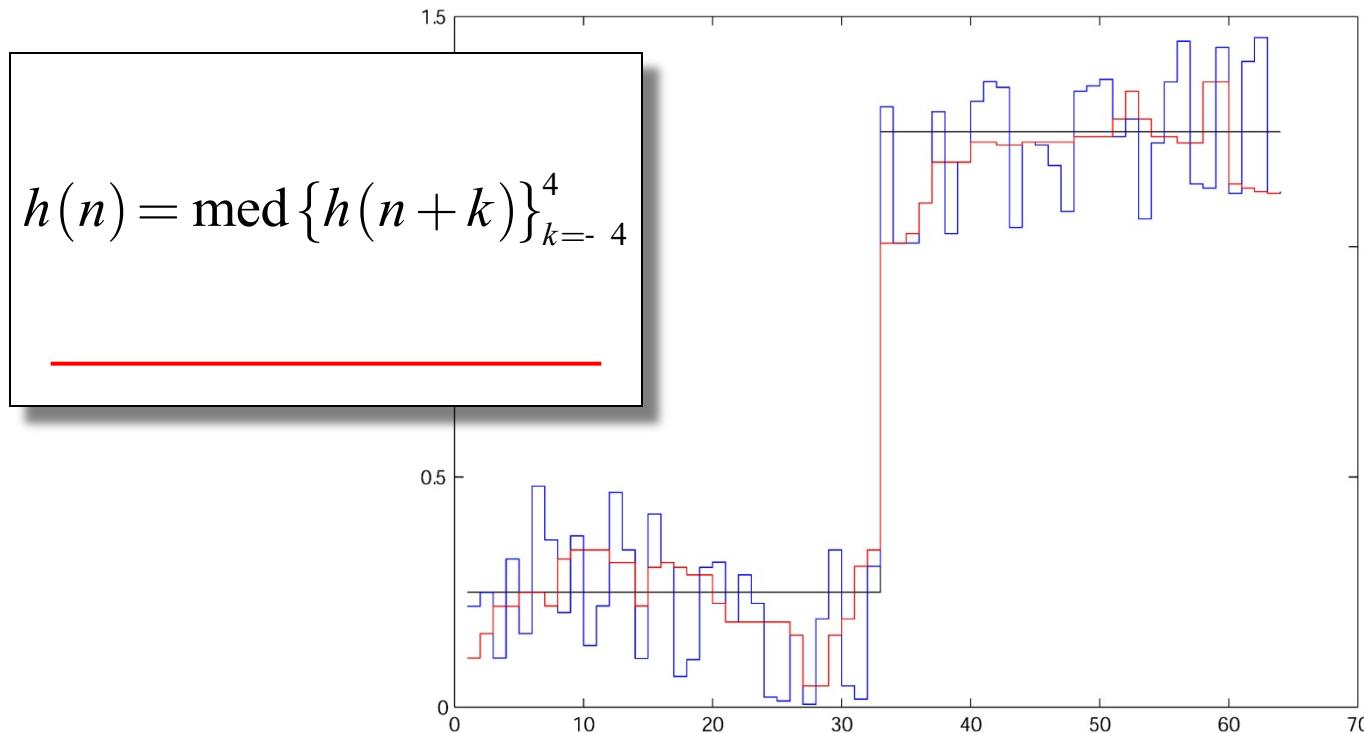
Blurred Noisy 1D Step Edge



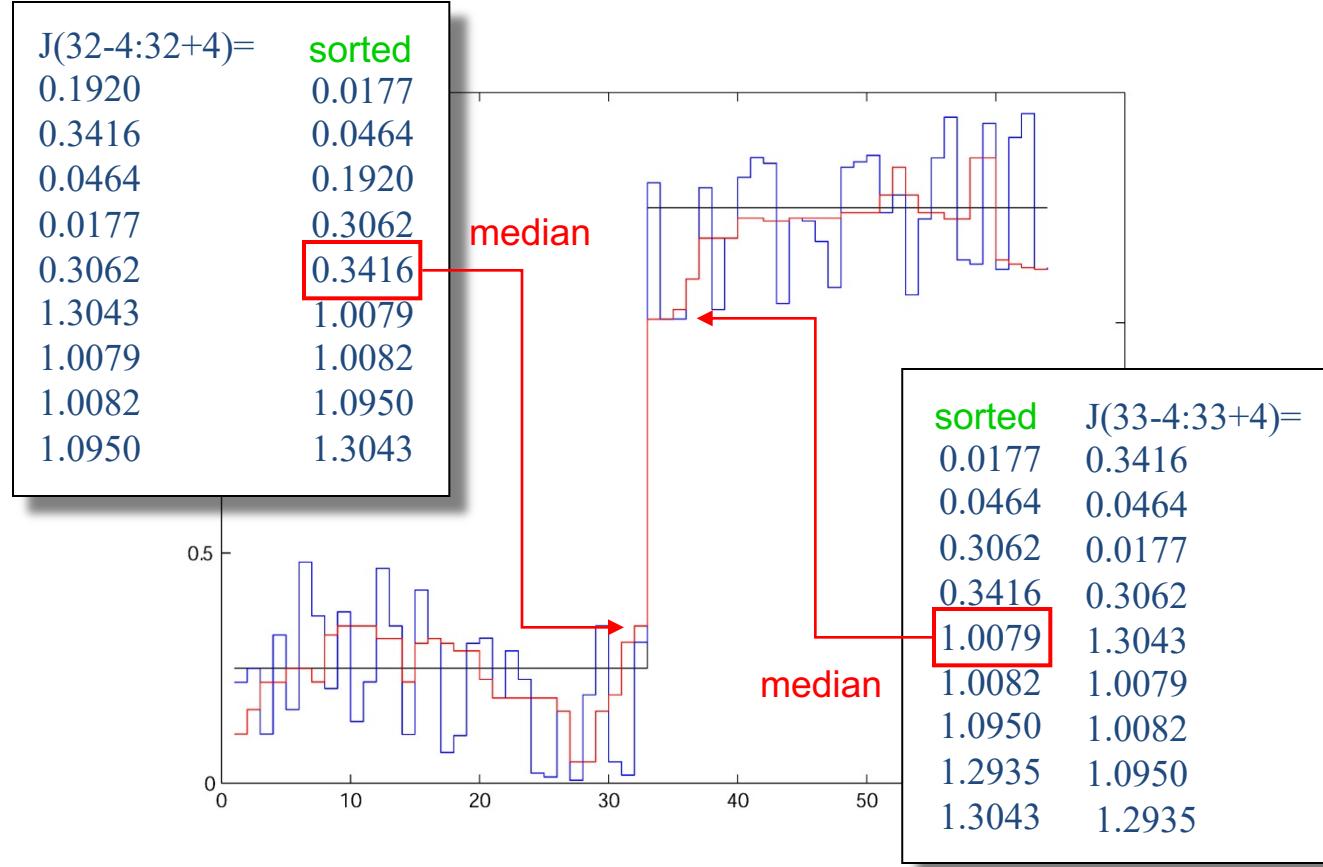
Blurred Noisy 1D Step Edge



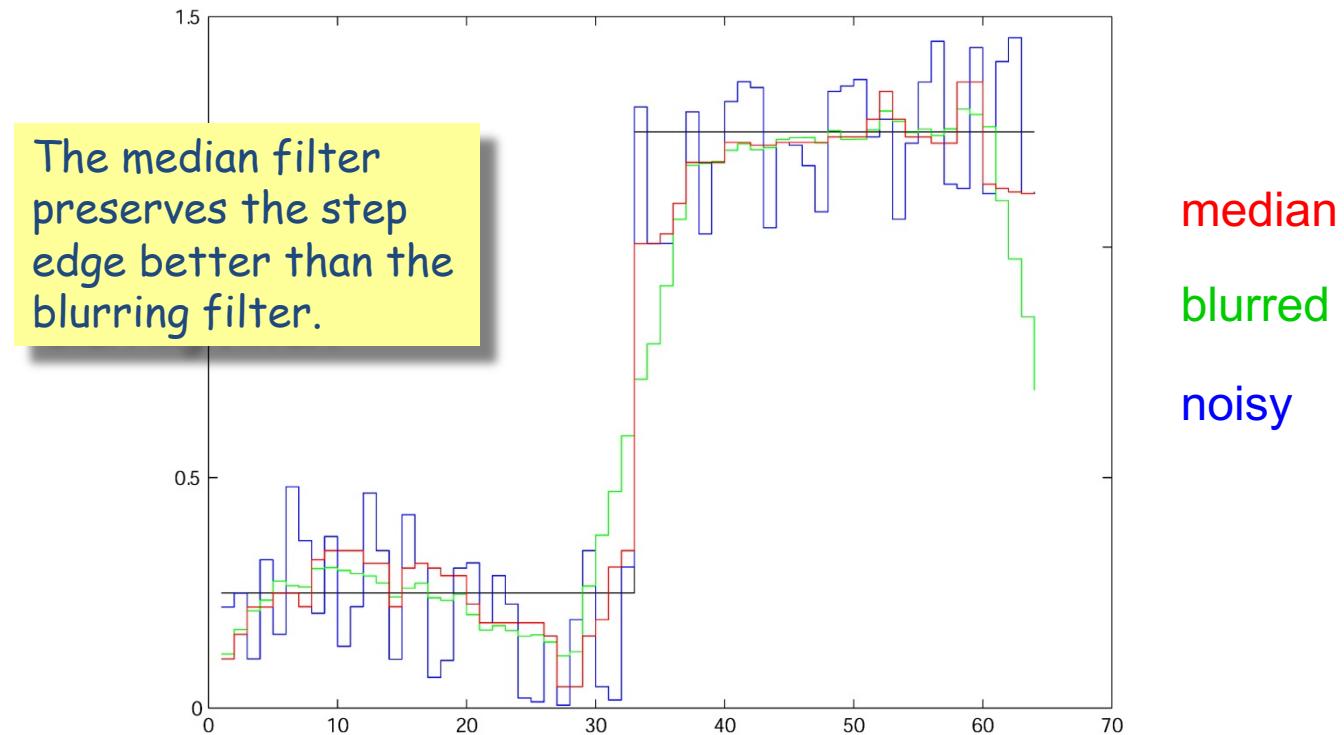
Median Filtered Noisy 1D Step Edge



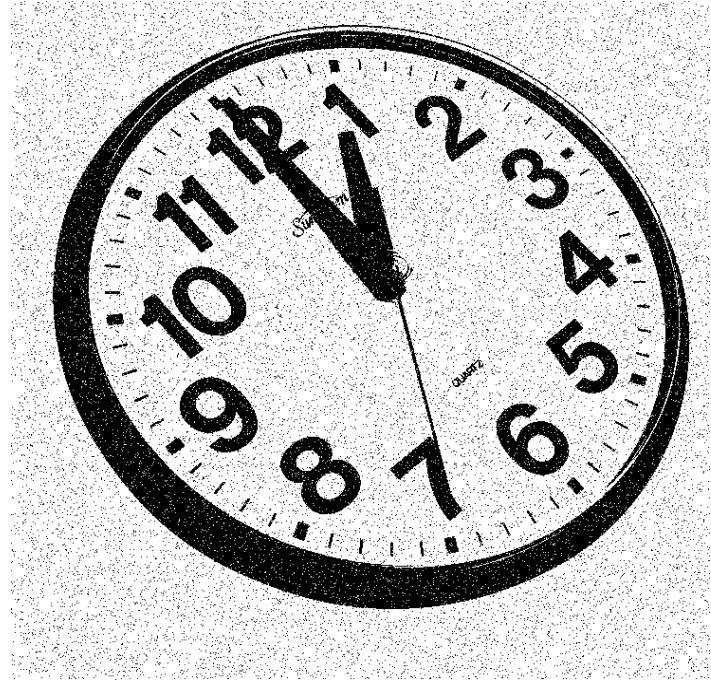
Median Filtered Noisy 1D Step Edge



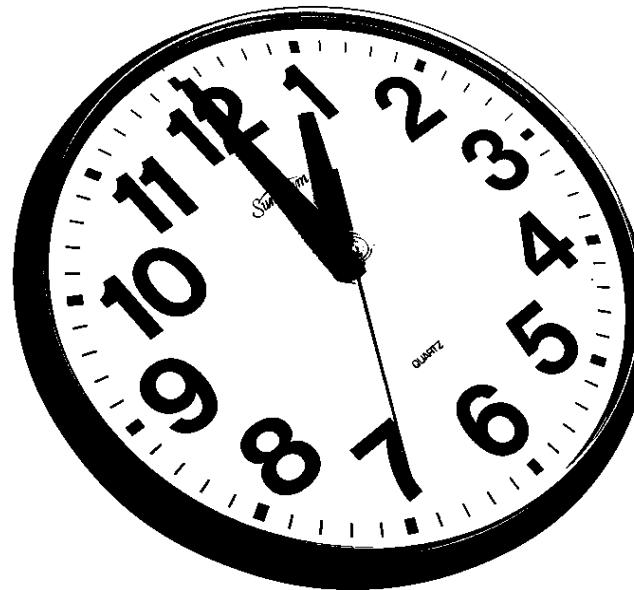
Median vs. Blurred



Median Filtering of Binary Images

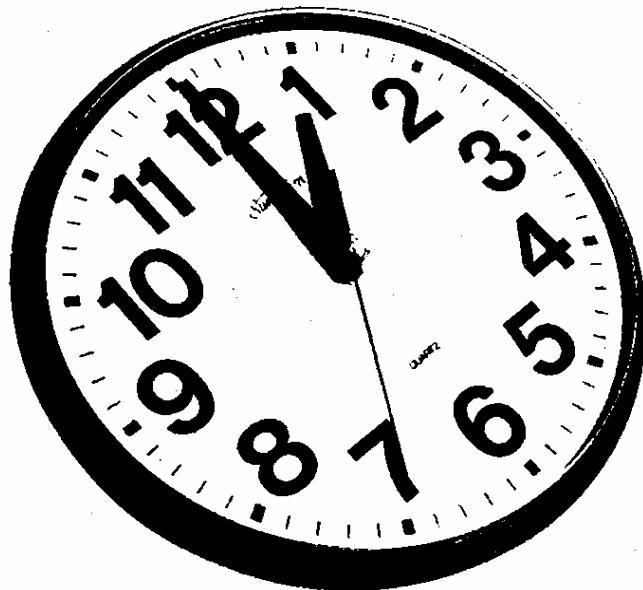


Noisy

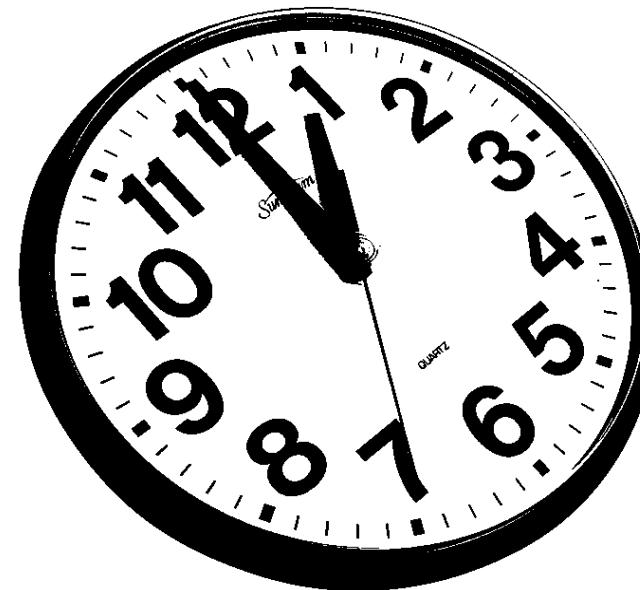


Original

Median Filtering of Binary Images

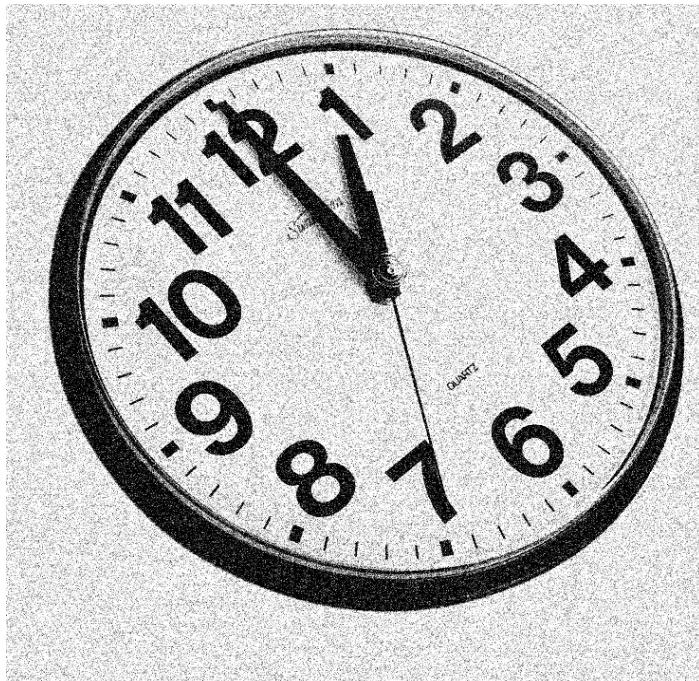


Median Filtered Noisy



Original

Filtering of Grayscale Images

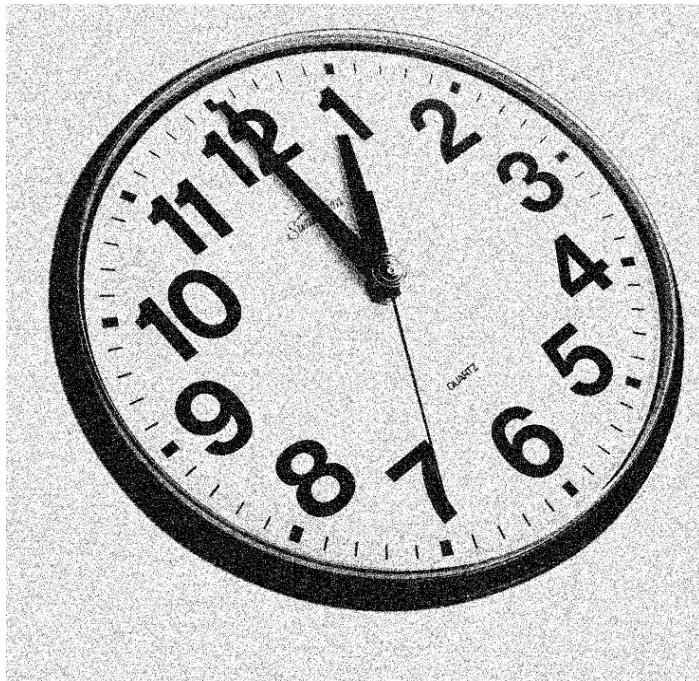


Noisy

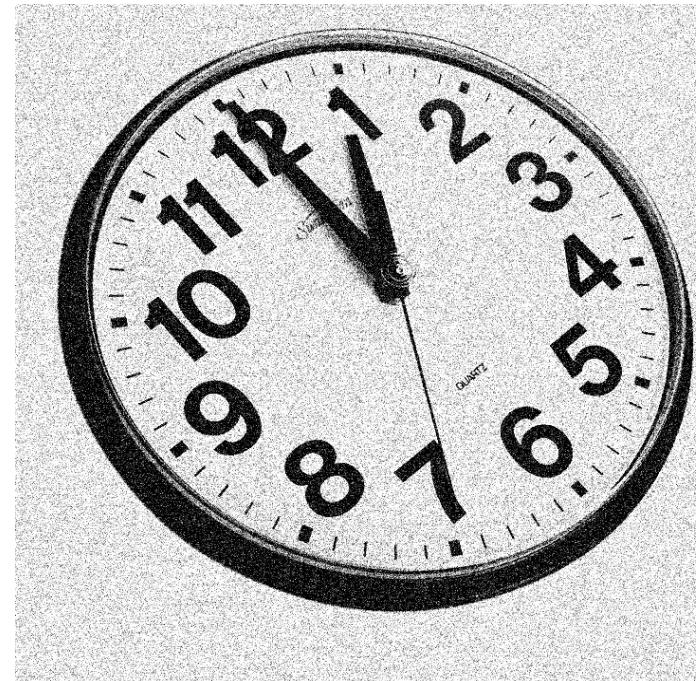


Original

Filtering of Grayscale Images

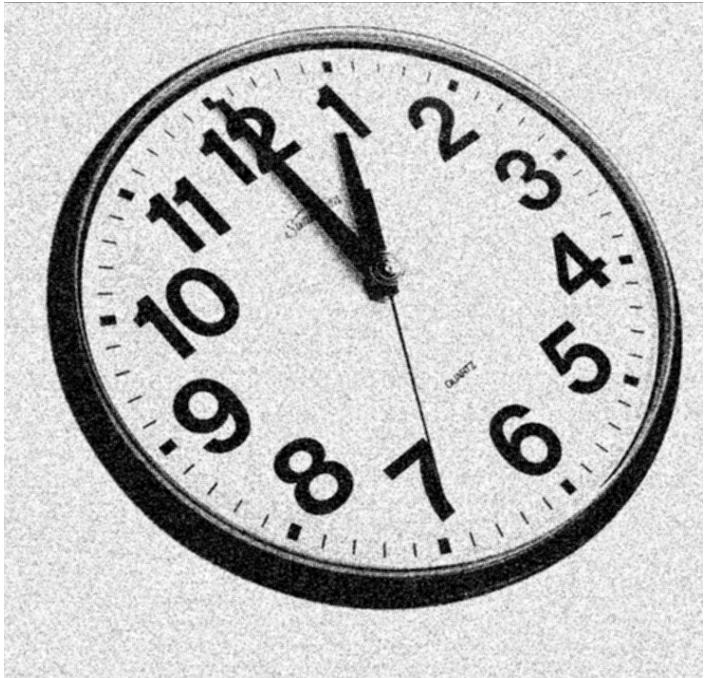


Noisy

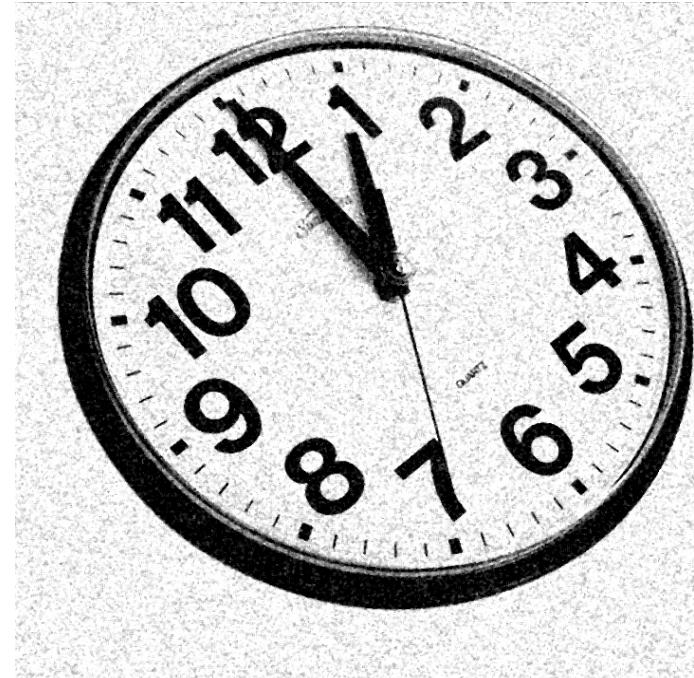


Noisy

Filtering of Grayscale Images

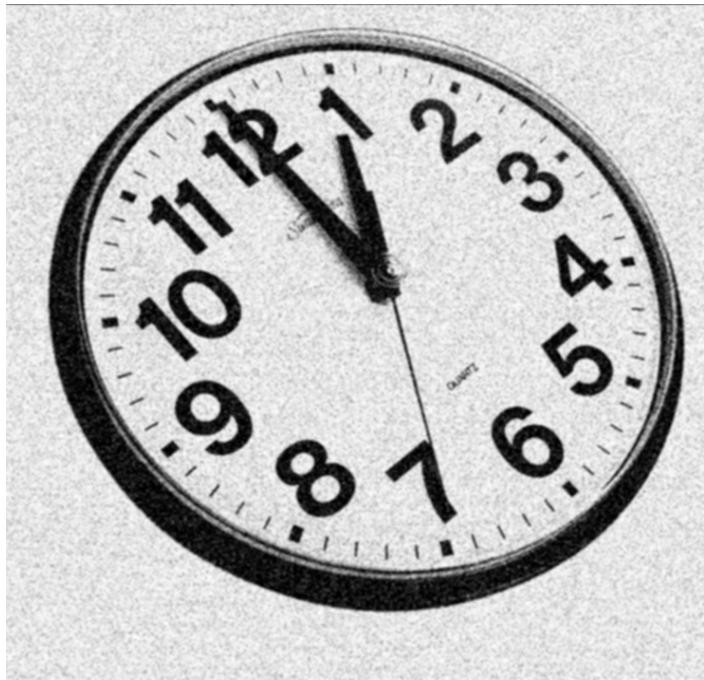


3x3-blur x 1

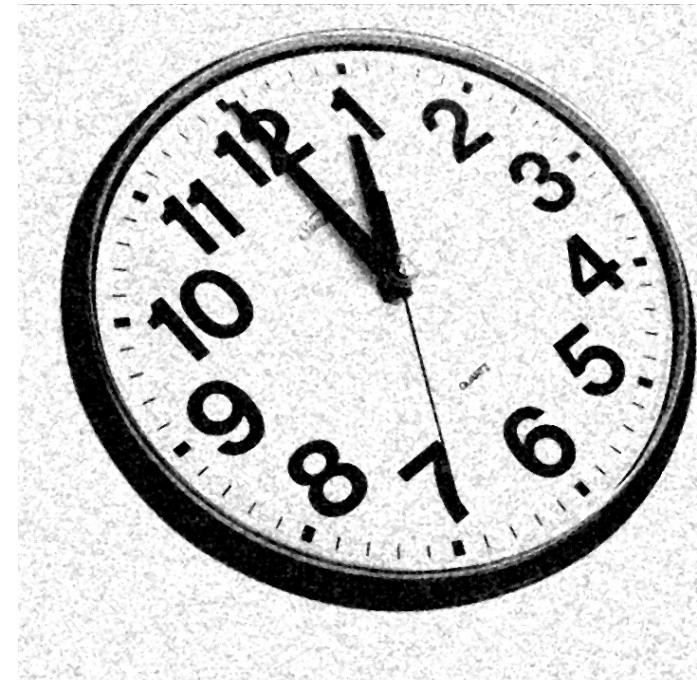


3x3-median x 1

Filtering of Grayscale Images

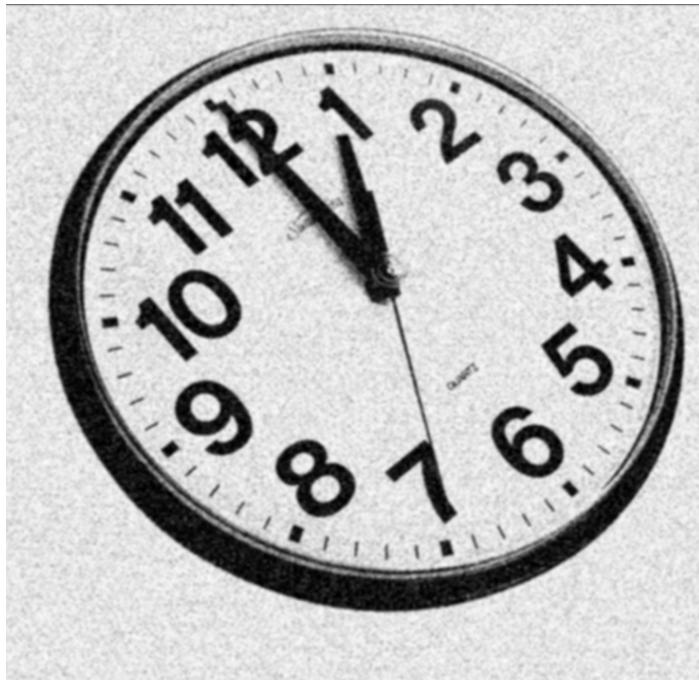


3x3-blur x 2

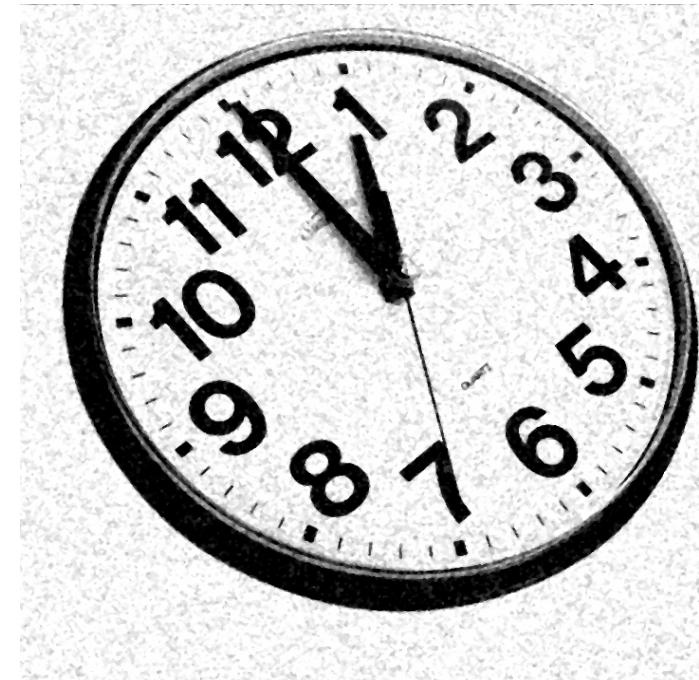


3x3-median x 2

Filtering of Grayscale Images

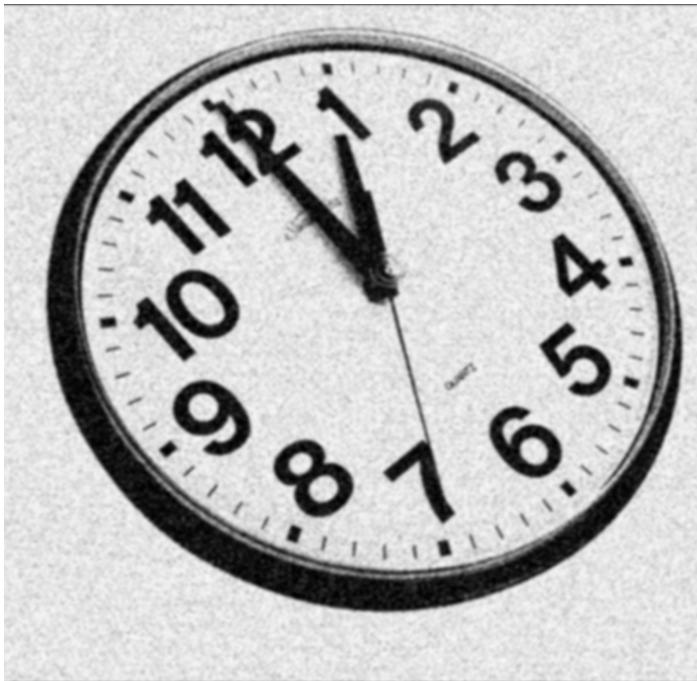


3x3-blur x 3

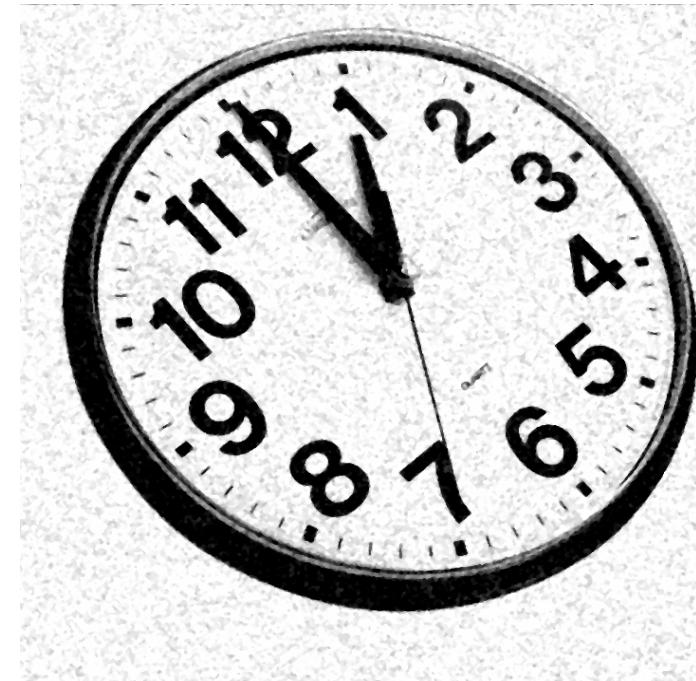


3x3-median x 3

Filtering of Grayscale Images

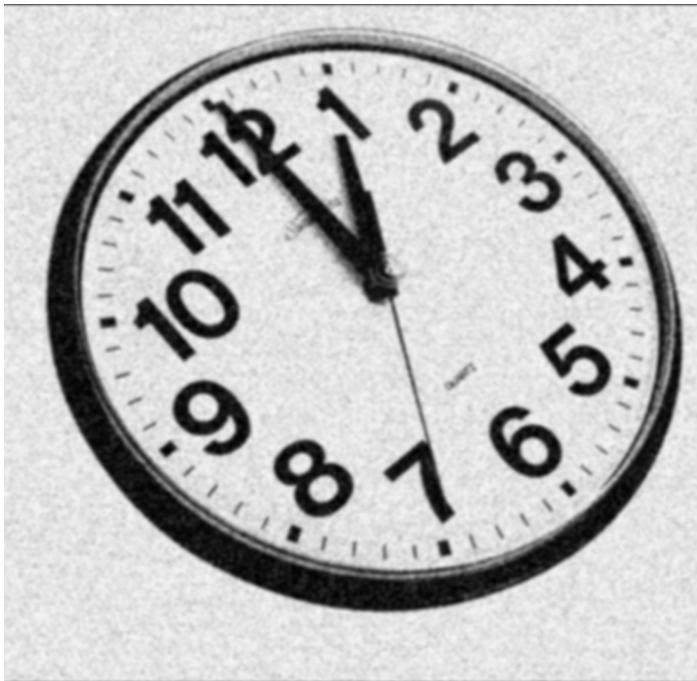


3x3-blur x 4

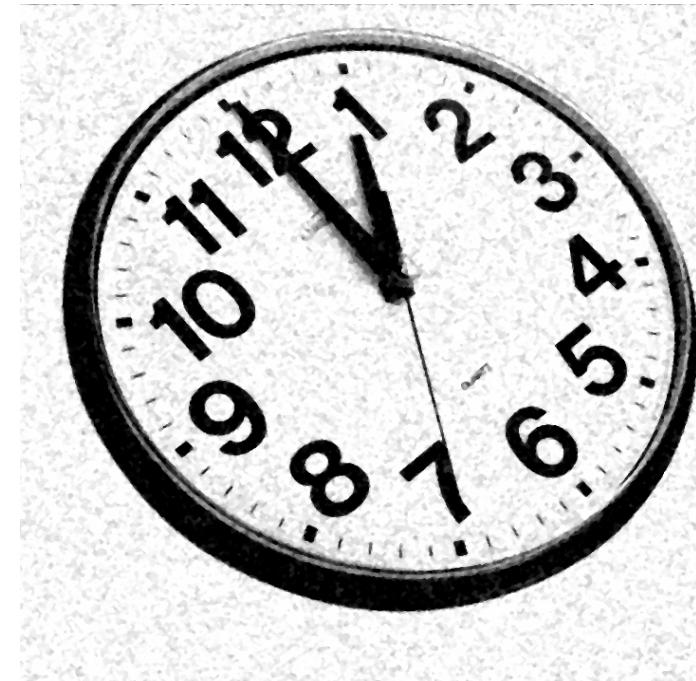


3x3-median x 4

Filtering of Grayscale Images



3x3-blur x 5

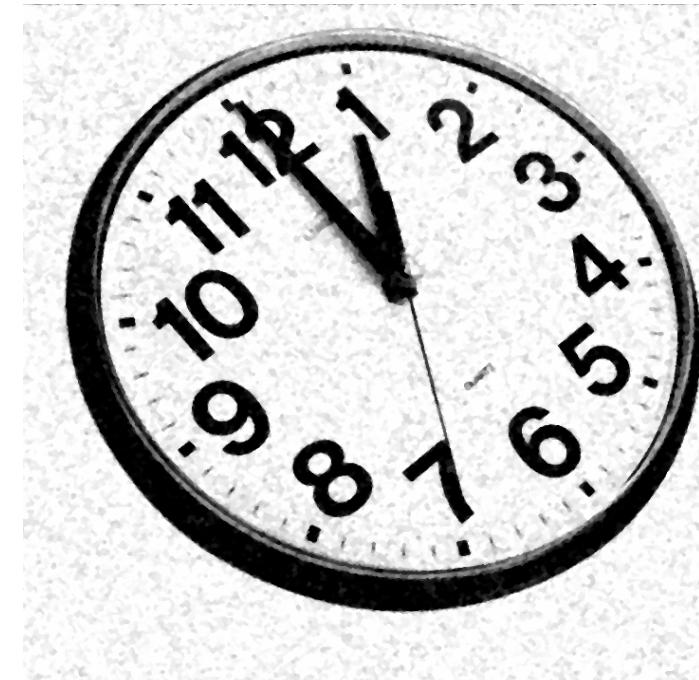


3x3-median x 5

Filtering of Grayscale Images



3x3-blur x 10



3x3-median x 10

Limit and Root Images

Fact: if you repeatedly filter an image with the same blurring filter or median filter, eventually the output does not change. That is, let

$$\mathbf{I}[*\mathbf{h}]^k \equiv (((\mathbf{I} * \mathbf{h}) * \mathbf{h}) \cdots * \mathbf{h}), \quad k \text{ times, and}$$

$$\mathbf{I}[\text{med } \mathbf{Z}]^k \equiv (((\mathbf{I} \text{ med } \mathbf{Z}) \text{ med } \mathbf{Z}) \cdots \text{med } \mathbf{Z}), \quad k \text{ times.}$$

Then

$$\lim_{k \rightarrow \infty} \mathbf{I}[*\mathbf{h}]^k = \mathbf{I}[*\mathbf{h}]^n = \mathbf{I}_0, \quad \text{and}$$

$$\lim_{k \rightarrow \infty} \mathbf{I}[\text{med } \mathbf{Z}]^k = \mathbf{I}[\text{med } \mathbf{Z}]^m = \mathbf{I}_r,$$

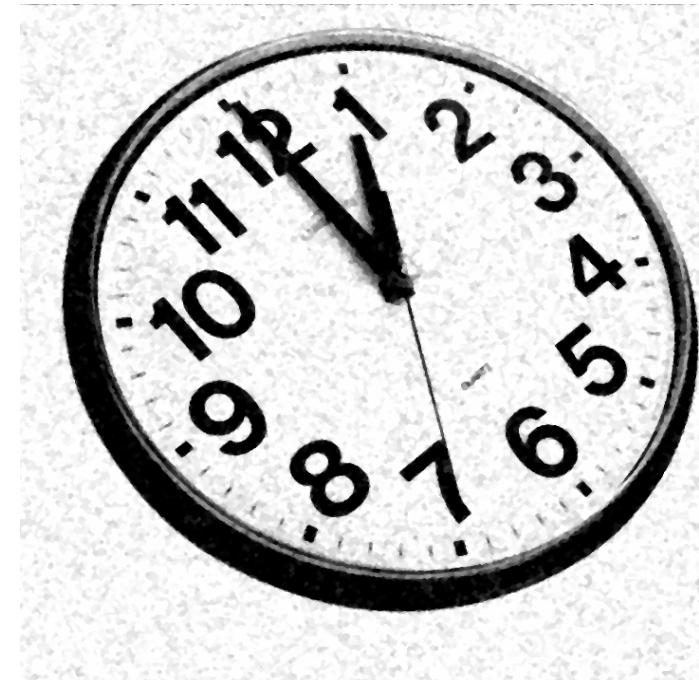
where n and m are integers ($< \infty$) , \mathbf{I}_0 is a single-valued image and \mathbf{I}_r is called the *median root* of \mathbf{I} .



Limit and Root Images



3x3-blur x 10

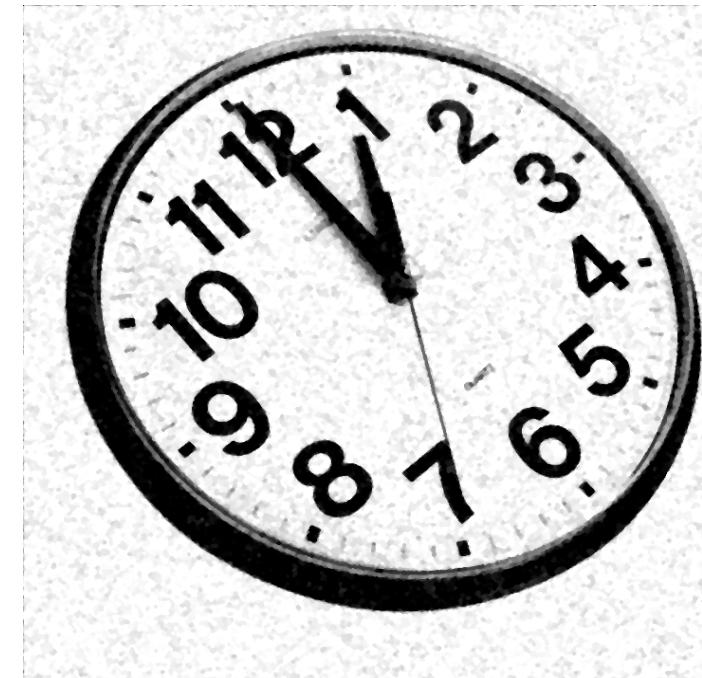


3x3-median x 10

Limit and Root Images



$3 \times 3\text{-blur } \times n \rightarrow \infty$



3x3-median root

Median Filter Algorithm in Matlab

```
function D = median_filt(I,SE,origy,origx)
[R,C] = size(I);      % assumes 1-band image
[SER,SEC] = size(SE); % SE < 0 not in nbhd

N = sum(sum(SE>=0)); % no. of pixels in nbhd
A = -ones(R+SER-1,C+SEC-1,N); % accumulator
n=1; % copy I into band n of A for nbhd pix n
for j = 1 : SER % neighborhood is def'd in SE
    for i = 1 : SEC
        if SE(j,i) >= 0 % then is a nbhd pixel
            A(j:(R+j-1),i:(C+i-1),n) = I;
            n=n+1; % next accumulator band
        end
    end
end
% pixel-wise median across the bands of A
A = shiftdim(median(shiftdim(A,2)),1);
D = A( origy:(R+origy-1) , origx:(C+origx-1) );
return;
```



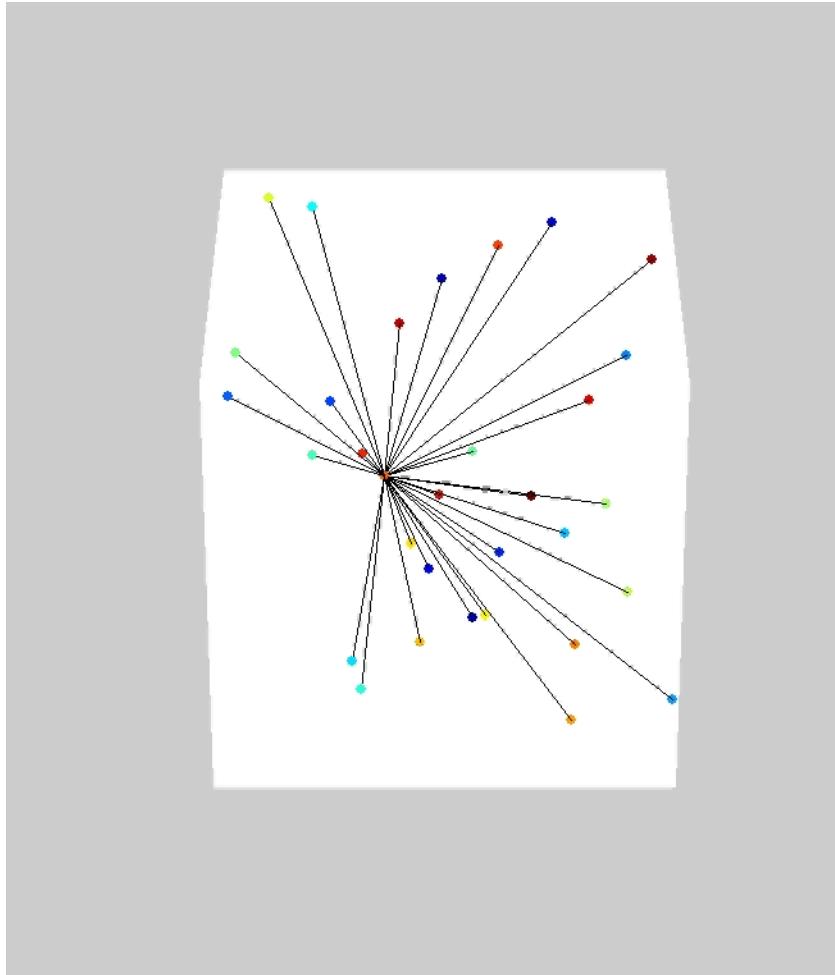
Vector Median Filter

A vector median filter selects from among a set of vectors, the one vector that is closest to all the others.

That is, if S is a set of vectors, in \mathbb{F}^n the median, $\bar{\mathbf{v}}$, is

$$\bar{\mathbf{v}} = \arg \min_{k \neq j} \left\{ \|\mathbf{v}_k - \mathbf{v}_j\| \mid \mathbf{v}_k, \mathbf{v}_j \in S \right\}.$$

($\trianglelefteq \mathbb{F}^n$ is an n-dimensional linear vector space over the field, \mathbb{F} .)



Color Median Filter



Jim Woodring – A Warm Shoulder



Sparse noise, 32% coverage in each band

Color Median Filter



3×3 color median filter applied once



3×3 color median filter applied twice

Color Median Filter

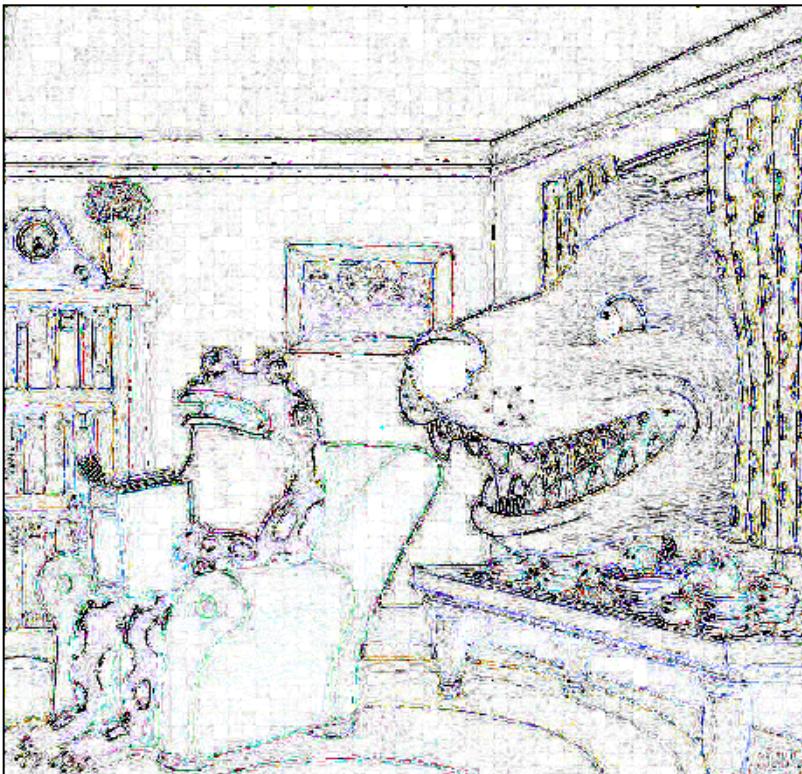


Sparse noise, 32% coverage in each band

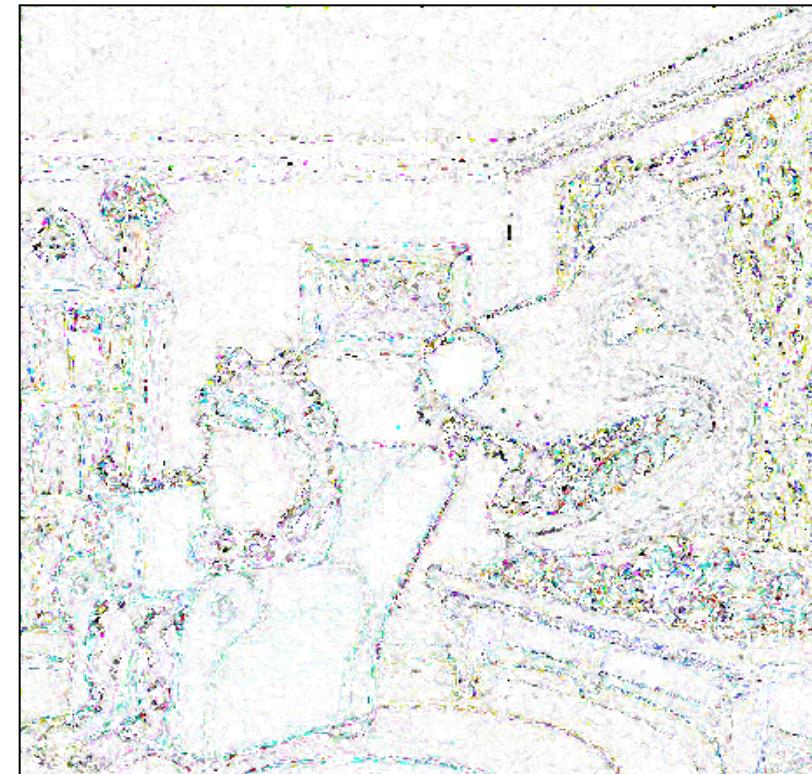


Jim Woodring – A Warm Shoulder

Color Median Filter



(3×3 CMF² of noisy) – original



(3×3 CMF² of noisy) – (3×3 CMF² of original)

Absolute differences
displayed as negatives
to enhance visibility

CMF vs. Standard Median on Individual Bands

A color median filter has to compute the distances between all the color vectors in the neighborhood of each pixel. That's expensive computationally.

Q: Why not simply take the 1-band median of each color band individually?

A: The result at a pixel could be a color that did not exist in the pixel's neighborhood in the input image. The result is not the median of the colors – it is the median of the intensities of each color band treated independently.

Q: Is that a problem?

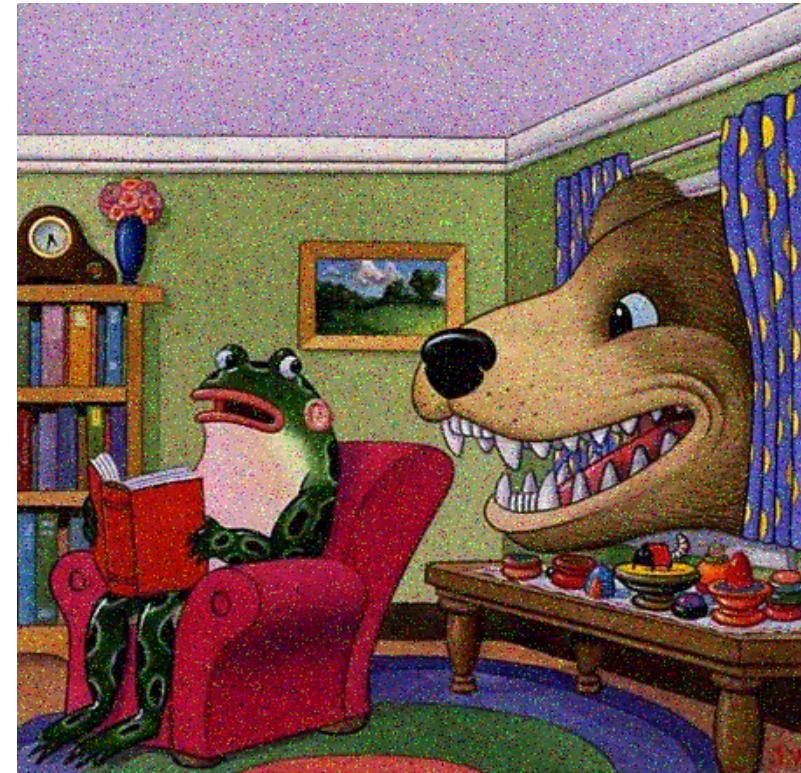
A: Maybe. Maybe not. It depends on the application. It may make little difference visually. If the colors need to be preserved, it could be problematic.



CMF vs. Standard Median on Individual Bands



Jim Woodring – A Warm Shoulder



Sparse noise, 32% coverage in each band

CMF vs. Standard Median on Individual Bands



3×3 color median filter applied once



3×3 color median filter applied twice

CMF vs. Standard Median on Individual Bands



3×3 median filter applied to each band once



3×3 median filter applied to each band twice