

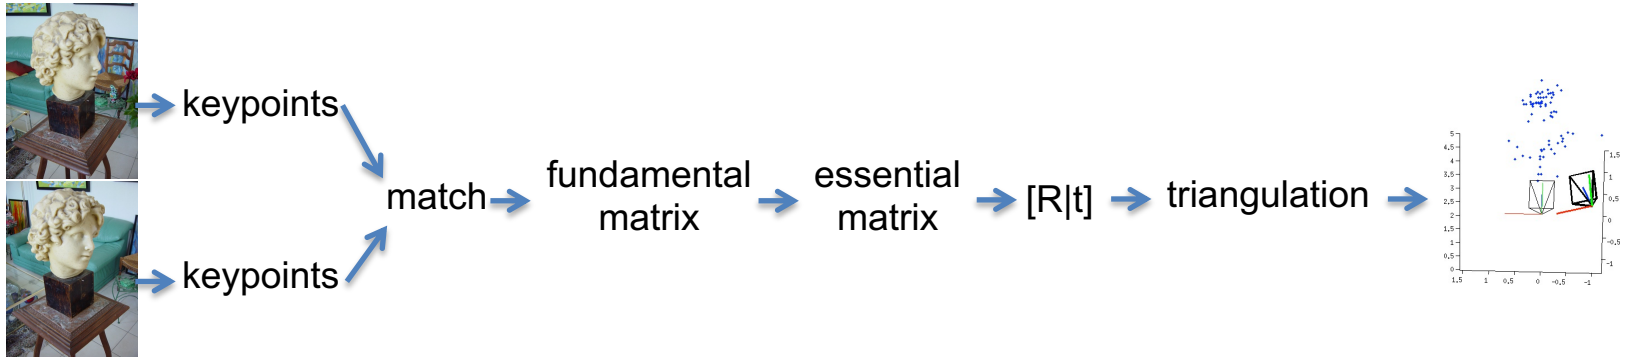


中国科学技术大学
University of Science and Technology of China

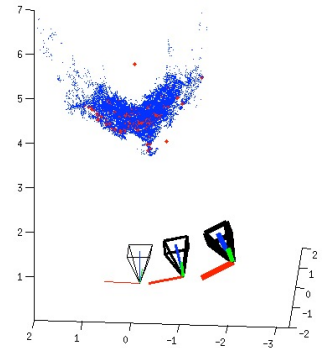





Bundle Adjustment

张举勇
中国科学技术大学

Two-view Reconstruction









Pipeline



Structure from Motion (SFM)

Multi-view Stereo (MVS)





Pipeline






Taught

Next

Merge Two Point Cloud



Merge Two Point Cloud



There can be only one $[\mathbf{R}_2 | \mathbf{t}_2]$

Merge Two Point Cloud

- From the 1st and 2nd images, we have




$$[\mathbf{R}_1 | \mathbf{t}_1] \text{ and } [\mathbf{R}_2 | \mathbf{t}_2]$$

- From the 2nd and 3rd images, we have







$$[\mathbf{R}_2 | \mathbf{t}_2] \text{ and } [\mathbf{R}_3 | \mathbf{t}_3]$$

- How to transform the coordinate system of the second point cloud to align with the first point cloud so that there is only one $[\mathbf{R}_2 | \mathbf{t}_2]$?

Merge Two Point Cloud




Oops



See From a Different Angle

Bundle Adjustment



Rethinking the SFM problem

- Input: Observed 2D image position

$$\tilde{\mathbf{x}}_1^1 \quad \tilde{\mathbf{x}}_1^2$$

$$\tilde{\mathbf{x}}_2^1 \quad \tilde{\mathbf{x}}_2^2 \quad \tilde{\mathbf{x}}_2^3$$

- Output: $\tilde{\mathbf{x}}_3^1 \quad \tilde{\mathbf{x}}_3^3$

Unknown Camera Parameters (with some guess)

$$[\mathbf{R}_1 | \mathbf{t}_1], [\mathbf{R}_2 | \mathbf{t}_2], [\mathbf{R}_3 | \mathbf{t}_3]$$

Unknown Point 3D coordinate (with some guess)

$$\mathbf{X}^1, \mathbf{X}^2, \mathbf{X}^3, \dots$$

Bundle Adjustment

A valid solution $[\mathbf{R}_1|\mathbf{t}_1], [\mathbf{R}_2|\mathbf{t}_2], [\mathbf{R}_3|\mathbf{t}_3]$ and $\mathbf{X}^1, \mathbf{X}^2, \mathbf{X}^3, \dots$ must let

$$\text{Re-projection} \left\{ \begin{array}{lll} \mathbf{x}_1^1 = \mathbf{K}[\mathbf{R}_1|\mathbf{t}_1]\mathbf{X}^1 & \mathbf{x}_1^2 = \mathbf{K}[\mathbf{R}_1|\mathbf{t}_1]\mathbf{X}^2 & \\ \mathbf{x}_2^1 = \mathbf{K}[\mathbf{R}_2|\mathbf{t}_2]\mathbf{X}^1 & \mathbf{x}_2^2 = \mathbf{K}[\mathbf{R}_2|\mathbf{t}_2]\mathbf{X}^2 & \mathbf{x}_2^3 = \mathbf{K}[\mathbf{R}_2|\mathbf{t}_2]\mathbf{X}^3 \\ \mathbf{x}_3^1 = \mathbf{K}[\mathbf{R}_3|\mathbf{t}_3]\mathbf{X}^1 & & \mathbf{x}_3^3 = \mathbf{K}[\mathbf{R}_3|\mathbf{t}_3]\mathbf{X}^3 \end{array} \right.$$

=

$$\text{Observation} \left\{ \begin{array}{lll} \tilde{\mathbf{x}}_1^1 & \tilde{\mathbf{x}}_1^2 & \\ \tilde{\mathbf{x}}_2^1 & \tilde{\mathbf{x}}_2^2 & \tilde{\mathbf{x}}_2^3 \\ \tilde{\mathbf{x}}_3^1 & & \tilde{\mathbf{x}}_3^3 \end{array} \right.$$

Bundle Adjustment

A valid solution $[\mathbf{R}_1 | \mathbf{t}_1], [\mathbf{R}_2 | \mathbf{t}_2], [\mathbf{R}_3 | \mathbf{t}_3]$ and $\mathbf{X}^1, \mathbf{X}^2, \mathbf{X}^3, \dots$ must let the Re-projection close to the Observation, i.e. to minimize the reprojection error

$$\min \sum_i \sum_j \left(\tilde{\mathbf{x}}_i^j - \mathbf{K}[\mathbf{R}_i | \mathbf{t}_i] \mathbf{X}^j \right)^2$$

Solving This Optimization Problem

- Theory:

The Levenberg–Marquardt algorithm

http://en.wikipedia.org/wiki/Levenberg-Marquardt_algorithm

- Practice:

The Ceres-Solver from Google

<http://code.google.com/p/ceres-solver/>

Ceres-solver: A Nonlinear Least Squares Minimizer

Toy problem to solve $\min (10 - x)^2$

```
class SimpleCostFunction
: public ceres::SizedCostFunction<1 /* number of residuals */,
                                1 /* size of first parameter */> {
public:
    virtual ~SimpleCostFunction() {}
    virtual bool Evaluate(double const* const* parameters,
                          double* residuals,
                          double** jacobians) const {
        const double x = parameters[0][0];
        residuals[0] = 10 - x; // f(x) = 10 - x
        // Compute the Jacobian if asked for.
        if (jacobians != NULL && jacobians[0] != NULL) {
            jacobians[0][0] = -1;
        }
        return true;
    }
};
```

Ceres-solver: A Nonlinear Least Squares Minimizer

Toy problem to solve $\min (10 - x)^2$

```
int main(int argc, char** argv) {  
    double x = 5.0;  
    ceres::Problem problem;  
  
    // The problem object takes ownership of the newly allocated  
    // SimpleCostFunction and uses it to optimize the value of x.  
    problem.AddResidualBlock(new SimpleCostFunction, NULL, &x);  
  
    // Run the solver!  
    Solver::Options options;  
    options.max_num_iterations = 10;  
    options.linear_solver_type = ceres::DENSE_QR;  
    options.minimizer_progress_to_stdout = true;  
    Solver::Summary summary;  
    Solve(options, &problem, &summary);  
    std::cout << summary.BriefReport() << "\n";  
    std::cout << "x : 5.0 -> " << x << "\n";  
    return 0;  
}
```


Ceres-solver: A Nonlinear Least Squares Minimizer

Toy problem to solve $\min (10 - x)^2$

```
0: f: 1.250000e+01 d: 0.00e+00 g: 5.00e+00 h: 0.00e+00 rho: 0.00e+00 mu: 1.00e-04 li: 0
1: f: 1.249750e-07 d: 1.25e+01 g: 5.00e-04 h: 5.00e+00 rho: 1.00e+00 mu: 3.33e-05 li: 1
2: f: 1.388518e-16 d: 1.25e-07 g: 1.67e-08 h: 5.00e-04 rho: 1.00e+00 mu: 1.11e-05 li: 1
Ceres Solver Report: Iterations: 2, Initial cost: 1.250000e+01, \
Final cost: 1.388518e-16, Termination: PARAMETER_TOLERANCE.
x : 5 -> 10
```